

Chapter 2

Mathematical Foundation

supplementary slides to
Machine Learning Fundamentals
©Hui Jiang 2020
published by Cambridge University Press

August 2020



CAMBRIDGE
UNIVERSITY PRESS

Outline

1 Linear Algebra

2 Probability and Statistics

3 Information Theory

4 Mathematical Optimization

Vectors and Matrices

- a vector: a list of numbers arranged in order
 - an abstract way to represent objects in math
 - vectors are written in a column, e.g. $\mathbf{x} \in \mathbb{R}^n$ and $\mathbf{y} \in \mathbb{R}^m$

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{bmatrix}$$

- a matrix: a group of numbers arranged in a 2-d array, e.g.
 $\mathbf{A} \in \mathbb{R}^{m \times n}$

$$\mathbf{A} = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{bmatrix}$$

Matrix Multiplication (I)

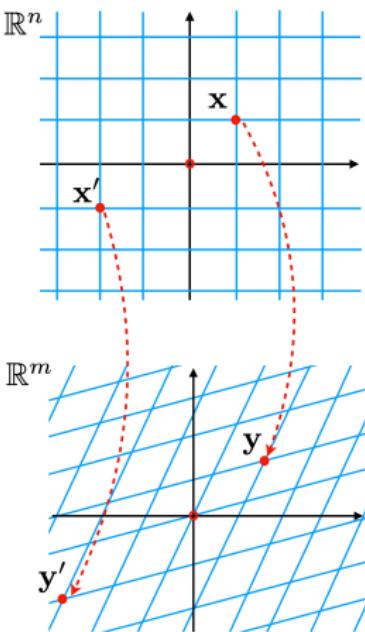
- a matrix: representing a particular way to move any point in one space to another
- matrix multiplication: the way to implement the above motion

$$\underbrace{\begin{bmatrix} \mathbf{y} \\ \vdots \\ y_i \\ \vdots \\ y_m \end{bmatrix}}_{\mathbf{y}} = \underbrace{\begin{bmatrix} a_{11} & a_{12} & \cdots & \cdots & a_{1n} \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ a_{i1} & \cdots & a_{ij} & \cdots & a_{in} \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ a_{m1} & a_{m2} & \cdots & \cdots & a_{mn} \end{bmatrix}}_A \underbrace{\begin{bmatrix} \mathbf{x} \\ \vdots \\ x_1 \\ \vdots \\ x_n \end{bmatrix}}_{\mathbf{x}}$$
$$y_i = \sum_{j=1}^n a_{ij} x_j \quad (\forall i = 1, 2, \dots, m)$$

$\mathbf{y} = \mathbf{Ax}$ ($\mathbf{A} \in \mathbb{R}^{m \times n}$) represents a mapping from a point \mathbf{x} in \mathbb{R}^n to another point \mathbf{y} in \mathbb{R}^m

Linear Transformation as Matrix Multiplication

- matrix multiplication: can only represent a **linear transformation**
 - a mapping from \mathbb{R}^n to \mathbb{R}^m is linear if and only if :
 - the origin in \mathbb{R}^n is mapped to the origin in \mathbb{R}^m
 - every straight line in \mathbb{R}^n is always mapped to a straight line (or a single point) in \mathbb{R}^m
 - nonlinear transformations: require other methods rather than matrix multiplication



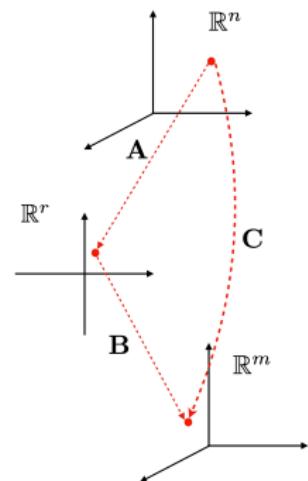
Matrix Multiplication (II)

- matrix multiplication between two matrices $\mathbf{A} \in \mathbb{R}^{m \times r}$ and $\mathbf{B} \in \mathbb{R}^{r \times n}$

$$\mathbf{C} = \mathbf{AB} \quad (\text{with } \mathbf{C} \in \mathbb{R}^{m \times n})$$

$$\begin{bmatrix} c_{11} & \cdots & c_{1n} \\ \vdots & & \vdots \\ c_{m1} & \cdots & c_{mn} \end{bmatrix}_C = \underset{\text{i-th row}}{\underset{\text{red}}{}} \begin{bmatrix} a_{11} & \cdots & a_{1r} \\ \vdots & & \vdots \\ a_{m1} & \cdots & a_{mr} \end{bmatrix}_A \underset{\text{j-th column}}{\underset{\text{red}}{}} \begin{bmatrix} b_{11} & \cdots & b_{1n} \\ \vdots & & \vdots \\ b_{r1} & \cdots & b_{rn} \end{bmatrix}_B$$

$$c_{ij} = \sum_{k=1}^r a_{ik} b_{kj} \quad \forall i = 1, 2, \dots, m \quad \forall j = 1, 2, \dots, n$$



- \mathbf{C} represents a composition of two linear transformations: \mathbf{A} and \mathbf{B}

Matrix Operations (I)

- **matrix transpose:** $\mathbf{A} \in \mathbb{R}^{m \times n} \rightarrow \mathbf{A}^T \in \mathbb{R}^{n \times m}$
 - a square matrix \mathbf{A} is symmetric iff $\mathbf{A} = \mathbf{A}^T$
 - $(\mathbf{A}^T)^T = \mathbf{A}$ $(\mathbf{AB})^T = \mathbf{B}^T \mathbf{A}^T$ $(\mathbf{A} \pm \mathbf{B})^T = \mathbf{A}^T \pm \mathbf{B}^T$
- **determinant:** $|\mathbf{A}| \in \mathbb{R}$ for a square matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$
- **inverse matrix \mathbf{A}^{-1} :** for an $n \times n$ square matrix \mathbf{A} if
 $\mathbf{A}^{-1}\mathbf{A} = \mathbf{AA}^{-1} = \mathbf{I}$
- **inner-product:** $\mathbf{w} \cdot \mathbf{x}$ ($\in \mathbb{R}$) of two vectors $\mathbf{w}, \mathbf{x} \in \mathbb{R}^n$
 - $\mathbf{w} \cdot \mathbf{x} \stackrel{\Delta}{=} \sum_{i=1}^n w_i x_i = \mathbf{w}^T \mathbf{x} = \mathbf{x}^T \mathbf{w}$
 - **norm** (L_2 norm) of vector \mathbf{w} : $\|\mathbf{w}\|^2 = \mathbf{w} \cdot \mathbf{w}$
- **trace** of an $n \times n$ square matrix: $\text{tr}(\mathbf{A}) = \sum_{i=1}^n a_{ii}$

Matrix Operations (II)

- **eigenvalues** and **eigenvectors** of an $n \times n$ square matrix \mathbf{A} :
 $\mathbf{A} \mathbf{u} = \lambda \mathbf{u}$ with $\lambda \in \mathbb{R}$ and non-zero $\mathbf{u} \in \mathbb{R}^n$
- A symmetric matrix is **positive definite** (or **semi-definite**) if all eigenvalues are positive (or non-negative).
- If a square matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$ has n orthogonal eigenvectors \mathbf{u}_i ($i = 1, 2, \dots, n$): $\mathbf{A} \mathbf{u}_i = \lambda_i \mathbf{u}_i$ (normalized as $\|\mathbf{u}_i\| = 1$), then \mathbf{A} can be factorized as $\mathbf{A} = \mathbf{U} \Lambda \mathbf{U}^\top$.

$$\mathbf{A} = \underbrace{\begin{bmatrix} | & | & & | \\ \mathbf{u}_1 & \mathbf{u}_2 & \cdots & \mathbf{u}_n \\ | & | & & | \end{bmatrix}}_{\mathbf{U}} \underbrace{\begin{bmatrix} \lambda_1 & 0 & \cdots & 0 \\ 0 & \lambda_2 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \lambda_n \end{bmatrix}}_{\Lambda} \underbrace{\begin{bmatrix} | & | & & | \\ \mathbf{u}_1 & \mathbf{u}_2 & \cdots & \mathbf{u}_n \\ | & | & & | \end{bmatrix}^\top}_{\mathbf{U}^\top}$$

Matrix Calculus

If y is a function involving all elements of a vector \mathbf{x} (or a matrix \mathbf{A}), $\frac{\partial y}{\partial \mathbf{x}}$ (or $\frac{\partial y}{\partial \mathbf{A}}$) is a vector (or a matrix) of the same size.

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \quad \text{and} \quad \mathbf{A} = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{bmatrix}$$

$$\frac{\partial y}{\partial \mathbf{x}} \triangleq \begin{bmatrix} \frac{\partial y}{\partial x_1} \\ \frac{\partial y}{\partial x_2} \\ \vdots \\ \frac{\partial y}{\partial x_n} \end{bmatrix} \quad \text{and} \quad \frac{\partial y}{\partial \mathbf{A}} \triangleq \begin{bmatrix} \frac{\partial y}{\partial a_{11}} & \frac{\partial y}{\partial a_{12}} & \cdots & \frac{\partial y}{\partial a_{1n}} \\ \frac{\partial y}{\partial a_{21}} & \frac{\partial y}{\partial a_{22}} & \cdots & \frac{\partial y}{\partial a_{2n}} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial y}{\partial a_{m1}} & \frac{\partial y}{\partial a_{m2}} & \cdots & \frac{\partial y}{\partial a_{mn}} \end{bmatrix}$$

Matrix Calculus Formula for Machine Learning (I)

$$\frac{\partial}{\partial \mathbf{x}} (\mathbf{x}^\top \mathbf{x}) = 2\mathbf{x}$$

$$\frac{\partial}{\partial \mathbf{x}} (\mathbf{x}^\top \mathbf{y}) = \mathbf{y}$$

$$\frac{\partial}{\partial \mathbf{x}} (\mathbf{x}^\top \mathbf{A} \mathbf{x}) = \mathbf{A}\mathbf{x} + \mathbf{A}^\top \mathbf{x}$$

$$\frac{\partial}{\partial \mathbf{x}} (\mathbf{x}^\top \mathbf{A} \mathbf{x}) = 2\mathbf{A}\mathbf{x} \quad (\text{symmetric } \mathbf{A})$$

$$\frac{\partial}{\partial \mathbf{A}} (\mathbf{x}^\top \mathbf{A} \mathbf{y}) = \mathbf{x}\mathbf{y}^\top$$

Matrix calculus formula for machine learning (II)

$$\frac{\partial}{\partial \mathbf{A}} \left(\mathbf{x}^\top \mathbf{A}^{-1} \mathbf{y} \right) = -(\mathbf{A}^\top)^{-1} \mathbf{x} \mathbf{y}^\top (\mathbf{A}^\top)^{-1} \quad (\text{square } \mathbf{A})$$

$$\frac{\partial}{\partial \mathbf{A}} \left(\ln |\mathbf{A}| \right) = (\mathbf{A}^{-1})^\top = (\mathbf{A}^\top)^{-1} \quad (\text{square } \mathbf{A})$$

$$\frac{\partial}{\partial \mathbf{A}} \left(\text{tr}(\mathbf{A}) \right) = \mathbf{I} \quad (\text{square } \mathbf{A})$$

Random Variables

- probability $\Pr(A)$: a real number between 0 and 1, indicating how likely a random event A is to occur
- **random variables**: taking different values in different probabilities, e.g. X, Y, Z, \dots
 - discrete random variables (RVs)
 - continuous random variables (RVs)
- a random variable is fully specified by two ingredients:
 - its domain: the set of all possible values
 - its probability distribution: how likely it takes each value
- a probability function is used to characterize how likely a random variable may take each value in the domain

Probability Functions

- **probability mass functions (p.m.f)**
for discrete random variables
 - $p(x) \triangleq \Pr(X = x)$ in the domain
 $\forall x \in \{x_1, x_2, \dots\}$
 - sum-to-one constraint:
 $\sum_x p(x) = 1$

- **probability density functions (p.d.f)**
for continuous random variables
 - define $p(x)$ to ensure
 $\Pr(a \leq X \leq b) = \int_a^b p(x) dx$
 - sum-to-one constraint:
 $\int_{-\infty}^{+\infty} p(x) dx = 1$

x	x_1	x_2	x_3	x_4
$p(x)$	0.4	0.3	0.2	0.1

Table: an example of p.m.f

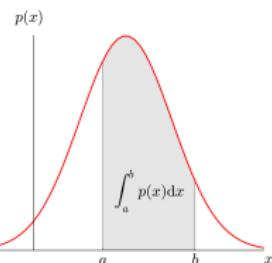


Figure: a p.d.f. in $(-\infty, \infty)$

Expectation: Mean, Variance and Moments

- **expectation** of any function of a random variable $f(X)$ is defined as

$$\mathbb{E}[f(X)] = \int_{-\infty}^{+\infty} f(x) p(x) dx \quad \text{or} \quad \mathbb{E}[f(X)] = \sum_x f(x) p(x)$$

- **mean** of a random variable X : $\mathbb{E}[X]$
 - the mean indicates the center of the distribution
- **variance** of X : $\text{var}(X) = \mathbb{E}[(X - \mathbb{E}[X])^2]$
 - the variance tells the average deviation from the center
- r -th **moment** of X : $\mathbb{E}[X^r]$ ($\forall r \in \mathbb{N}$)
- show: $\text{var}(X) = \mathbb{E}[X^2] - (\mathbb{E}[X])^2$

Joint Distributions

joint distributions of multiple random variables are multivariate probability functions defined in the product space of their domains

- joint p.m.f. for two discrete RVs

$$p(x, y) \triangleq \Pr(X = x, Y = y)$$

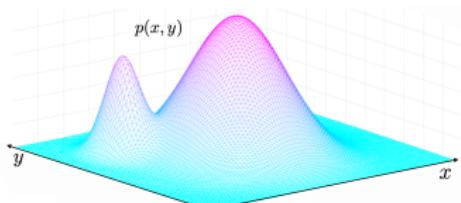
- $\forall x \in \{x_1, \dots\}, y \in \{y_1, \dots\}$
- $\sum_x \sum_y p(x, y) = 1$

- joint p.d.f. for two continuous RVs

$$\Pr((x, y) \in \Omega) = \int \int_{\Omega} p(x, y) dx dy$$

- $\int_{-\infty}^{+\infty} \int_{-\infty}^{+\infty} p(x, y) dx dy = 1$

$y \setminus x$	x_1	x_2	x_3
y_1	0.03	0.24	0.17
y_2	0.23	0.11	0.22



Marginal and Conditional Distributions

- a joint distribution fully specifies all random variables
- **marginalization** (a.k.a. *the rule of sum* in probability)
 - joint distribution → marginal distribution

$$p(x) = \int_{-\infty}^{+\infty} p(x, y) dy \quad \text{or} \quad p(x) = \sum_y p(x, y)$$

- **conditional distribution:** $p(x | y) \stackrel{\Delta}{=} \frac{p(x,y)}{p(y)}$
- *the general product rule* in probability:

$$p(x_1, x_2, x_3, \dots) = p(x_1) p(x_2|x_1) p(x_3|x_1, x_2) \dots$$

- conditional expectation: $\mathbb{E}_X [f(X) | Y = y_0]$
- conditional mean: $\mathbb{E}_X [X | Y = y_0]$
- covariance: $\text{cov}(X, Y) = \mathbb{E}[(X - \mathbb{E}[X])(Y - \mathbb{E}[Y])]$
- independence $\iff p(x, y) = p(x)p(y) \ (\forall x, y)$
- uncorrelatedness $\iff \text{cov}(X, Y) = 0$

Probability Distributions of Random Vectors

- random vector: whose elements are all random variables

$$p\left(\underbrace{x_1, x_2, x_3}_{\mathbf{x}}, \underbrace{y_1, y_2, y_3, y_4}_{\mathbf{y}}\right) = p(\mathbf{x}, \mathbf{y})$$

- conditional distribution: $p(\mathbf{x} | \mathbf{y}) \triangleq \frac{p(\mathbf{x}, \mathbf{y})}{p(\mathbf{y})}$
- mean vector: $\mathbb{E}[\mathbf{x}] = \int \mathbf{x} p(\mathbf{x}) d\mathbf{x} = \int \int \mathbf{x} p(\mathbf{x}, \mathbf{y}) d\mathbf{x} d\mathbf{y}$ or $\mathbb{E}[\mathbf{x}] = \sum_{\mathbf{x}} \sum_{\mathbf{y}} \mathbf{x} p(\mathbf{x}, \mathbf{y})$
- covariance matrix: $\text{cov}(\mathbf{x}, \mathbf{y}) = \mathbb{E}[(\mathbf{x} - \mathbb{E}[\mathbf{x}])(\mathbf{y} - \mathbb{E}[\mathbf{y}])^\top]$
- the rule of sum: $p(\mathbf{x}) = \int p(\mathbf{x}, \mathbf{y}) d\mathbf{y}$ or $p(\mathbf{x}) = \sum_{\mathbf{y}} p(\mathbf{x}, \mathbf{y})$
- the general product rule: $p(\mathbf{x}, \mathbf{y}, \mathbf{z}) = p(\mathbf{x}) p(\mathbf{y}|\mathbf{x}) p(\mathbf{z}|\mathbf{x}, \mathbf{y})$

Common Probability Distributions

- Binomial Distribution
- Multinomial Distribution
- Beta Distribution
- Dirichlet Distribution
- Univariate Gaussian (Normal) Distribution
- Multivariate Gaussian Distribution
- Poisson Distribution
- Uniform Distribution
- Gamma Distribution
- inverse-Wishart Distribution
- von Mises–Fisher Distribution

Binomial Distribution

$$B(r \mid N, p) \stackrel{\Delta}{=} \Pr(X = r) = \frac{N!}{r!(N-r)!} p^r (1-p)^{N-r}$$

- parameters: $N \in \mathbb{N}$ and $p \in [0, 1]$
- support: the domain of the random variable is $r \in \{0, 1, \dots, N\}$
- mean and variance:

$$\mathbb{E}[X] = Np$$

$$\text{var}(X) = Np(1-p)$$

- sum-to-one constraint:

$$\sum_{r=0}^N B(r \mid N, p) = 1$$

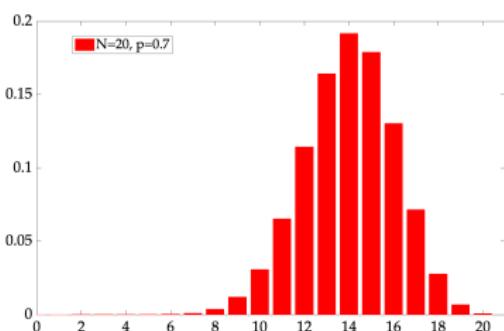


Figure: binomial distributions are **unimodal**

Multinomial Distribution

$$\begin{aligned}\text{Mult}(r_1, r_2, \dots, r_m \mid N, p_1, p_2, \dots, p_m) &\stackrel{\Delta}{=} \Pr(X_1 = r_1, \dots, X_m = r_m) \\ &= \frac{N!}{r_1! r_2! \cdots r_m!} p_1^{r_1} p_2^{r_2} \cdots p_m^{r_m}\end{aligned}$$

- multivariate extension of binomial distribution
- parameters: $N \in \mathbb{N}$; $0 \leq p_i \leq 1$ ($\forall i$) and $\sum_{i=1}^m p_i = 1$
- support (the domain of m random variables):
 $r_i \in \{0, 1, \dots, N\}$ ($\forall i = 1, \dots, m$) and $\sum_{i=1}^m r_i = N$
- means, variances and covariances:
 $\mathbb{E}[X_i] = Np_i$ and $\text{var}(X_i) = Np_i(1 - p_i)$ ($\forall i$)
 $\text{cov}(X_i, X_j) = -Np_i p_j$ ($\forall i, j$)
- sum-to-one: $\sum_{r_1 \dots r_m} \text{Mult}(r_1, \dots, r_m \mid N, p_1, \dots, p_m) = 1$

Beta Distribution

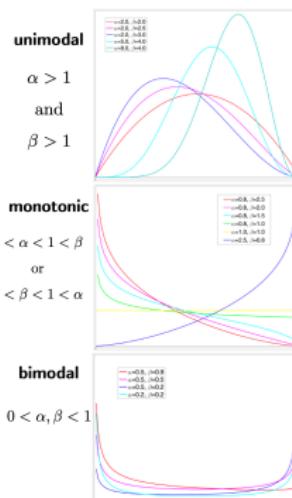
$$\text{Beta}(x | \alpha, \beta) = \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} x^{\alpha-1} (1-x)^{\beta-1}$$

- parameters: $\alpha > 0$ and $\beta > 0$
- support: $x \in \mathbb{R}$ and $0 \leq x \leq 1$
- mean and variance:

$$\mathbb{E}[X] = \frac{\alpha}{\alpha + \beta}$$

$$\text{var}(X) = \frac{\alpha\beta}{(\alpha + \beta)^2(\alpha + \beta + 1)}$$

- sum-to-one: $\int_0^1 \text{beta}(x | \alpha, \beta) dx = 1$



Dirichlet Distribution (I)

$$\text{Dir}(p_1, p_2, \dots, p_m \mid r_1, r_2, \dots, r_m) = \frac{\Gamma(r_1 + \dots + r_m)}{\Gamma(r_1) \dots \Gamma(r_m)} p_1^{r_1-1} p_2^{r_2-1} \dots p_m^{r_m-1}$$

- parameters: $r_i > 0$ ($\forall i = 1, \dots, m$)
- support is an m -dimensional simplex:
 $0 < p_i < 1$ ($\forall i = 1, \dots, m$) and $\sum_{i=1}^m p_i = 1$
- means, variances and covariances:
 $\mathbb{E}[X_i] = \frac{r_i}{r_0}$ $\text{var}(X_i) = \frac{r_i(r_0-r_i)}{r_0^2(r_0+1)}$
 $\text{cov}(X_i, X_j) = -\frac{r_i r_j}{r_0^2(r_0+1)}$, where $r_0 = \sum_{i=1}^m r_i$
- sum-to-one inside the simplex:
 $\int \dots \int_{p_1 \dots p_m} \text{Dir}(p_1, p_2, \dots, p_m \mid r_1, r_2, \dots, r_m) dp_1 \dots dp_m = 1$

Dirichlet Distribution (II)

- multivariate extension of beta distribution
- conjugate with multinomial distribution
- Dirichlet distributions ($m = 3$) in the 3-D simplex:

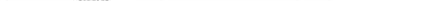
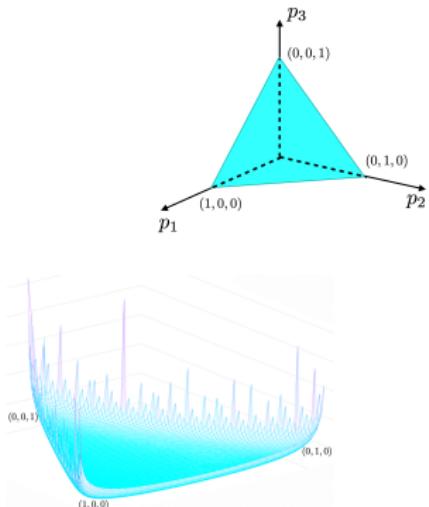
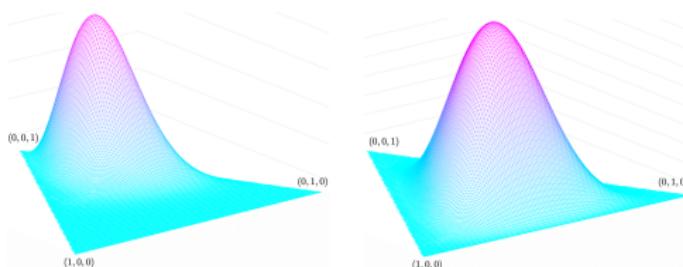


Figure: i) regular: $r_1 = 2.0, r_2 = 4.0, r_3 = 10.0$;
ii) symmetric: $r_1 = r_2 = r_3 = 4.0$;
iii) sparse: $r_1 = 0.7, r_2 = 0.8, r_3 = 0.9$;

Univariate Gaussian Distribution

$$\mathcal{N}(x | \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

- parameters: $\mu \in \mathbb{R}$ and $\sigma^2 > 0$
- support: $x \in \mathbb{R}$
- mean and variance:

$$\mathbb{E}[X] = \mu \quad \text{and} \quad \text{var}(X) = \sigma^2$$

- sum-to-one: $\int_{-\infty}^{+\infty} \mathcal{N}(x | \mu, \sigma) dx = 1$
- bell-shaped **unimodal**

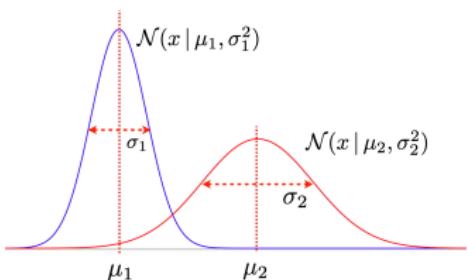


Figure: two univariate Gaussian distributions ($\sigma_2 > \sigma_1$).

Multivariate Gaussian Distribution

$$\mathcal{N}(\mathbf{x} | \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{\sqrt{(2\pi)^n |\boldsymbol{\Sigma}|}} e^{-\frac{(\mathbf{x}-\boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1} (\mathbf{x}-\boldsymbol{\mu})}{2}}$$

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$$

- parameters: $\boldsymbol{\mu} \in \mathbb{R}^n$ and $\boldsymbol{\Sigma} \in \mathbb{R}^{n \times n} \succ 0$
- support: $\mathbf{x} \in \mathbb{R}^n$
- mean vector and covariance matrix:

$$\mathbb{E}[\mathbf{x}] = \boldsymbol{\mu} \quad \text{and} \quad \text{cov}(\mathbf{x}, \mathbf{x}) = \boldsymbol{\Sigma}$$

- sum-to-one: $\int \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}, \boldsymbol{\Sigma}) d\mathbf{x} = 1$
- any marginal distribution or conditional distribution is Gaussian.
- multivariate unimodal

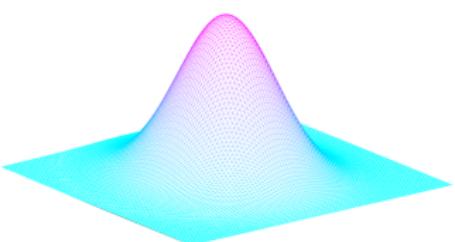


Figure: a multivariate Gaussian distribution in 2-dimensional space

Poisson Distribution

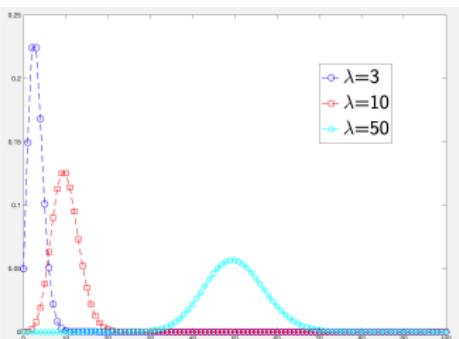
$$\text{Poisson}(n | \lambda) \triangleq \Pr(X = n) = \frac{e^{-\lambda} \cdot \lambda^n}{n!} \quad \forall n = 0, 1, 2 \dots$$

- parameter: $\lambda > 0$
- support: the domain of the random variable
 $n = 0, 1, 2, \dots$

- mean and variance:

$$\mathbb{E}[X] = \lambda \quad \text{and} \quad \text{var}(X) = \lambda$$

- sum-to-one: $\sum_{n=0}^{\infty} \text{Poisson}(n | \lambda) = 1$
- ideal distributions for count data



Transformation of Random Variables

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \xrightarrow{f} \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

If the transformation f is differentiable and invertible, then

$$p(\mathbf{y}) = |\mathbf{J}(\mathbf{y})| p(\mathbf{x}) = |\mathbf{J}(\mathbf{y})| p(f^{-1}(\mathbf{y}))$$

where $\mathbf{J}(\mathbf{y})$ denotes the Jacobian matrix of $\mathbf{x} = f^{-1}(\mathbf{y})$:

$$\mathbf{J}(\mathbf{y}) = \left[\frac{\partial x_i}{\partial y_j} \right]_{n \times n} = \begin{bmatrix} \frac{\partial x_1}{\partial y_1} & \dots & \frac{\partial x_1}{\partial y_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial x_n}{\partial y_1} & \dots & \frac{\partial x_n}{\partial y_n} \end{bmatrix}$$

Information Theory

- Shannon's **information** of an event: $I(A) = -\log_2 (\Pr(A))$
- **entropy** of a random variable: the total amount of uncertainty

$$H(X) = \mathbb{E}[-\log_2 \Pr(X = x)] = - \sum_x p(x) \log_2 p(x)$$

- the amount of information to resolve the random variable
- **joint entropy** of two random variables:

$$\begin{aligned} H(X, Y) &= \mathbb{E}_{X,Y} [-\log_2 \Pr(X = x, Y = y)] \\ &= - \sum_x \sum_y p(x, y) \log_2 p(x, y) \end{aligned}$$

- **conditional entropy**:

$$\begin{aligned} H(Y|X) &= \mathbb{E}_{X,Y} [-\log_2 \Pr(Y = y|X = x)] \\ &= - \sum_x \sum_y p(x, y) \log_2 p(y|x) \end{aligned}$$

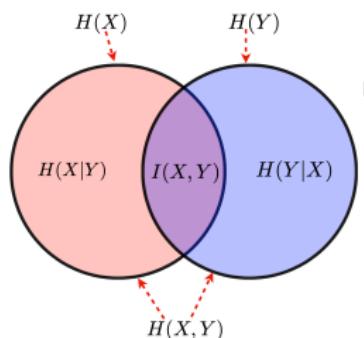
Mutual Information

■ mutual information:

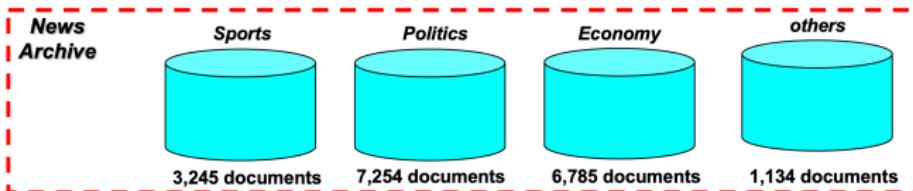
- entropy reduction of X after observing Y
- how much Y can tell about X

$$\begin{aligned} I(X, Y) &= H(X) - H(X|Y) \\ &= \sum_x \sum_y p(x, y) \log_2 \left(\frac{p(x, y)}{p(x)p(y)} \right) \end{aligned}$$

- $I(X, Y) = H(Y) - H(Y|X)$
 $= H(X) + H(Y) - H(X, Y)$
- symmetrical: $I(X, Y) = I(Y, X)$
- non-negative: $I(X, Y) \geq 0$
- $I(X, Y) = 0$ iff X and Y are independent



Example: Mutual Information for Keyword Selection



- text categorization: classify text documents into various topics
- text documents contain a large number of distinct words
- goal: use mutual information as a criterion to automatically filter out non-informative words
- for any word (e.g. *score*) and a topic (e.g. *sports*), define two binary random variables X and Y :
 - $X \in \{0, 1\}$: whether a document's topic is *sports* or not
 - $Y \in \{0, 1\}$: whether a document contains *score* or not
 - $I(X, Y) \implies$ relationship between *score* and *sports*

Example: Mutual Information for Keyword Selection

- count all training documents to estimate $p(X, Y)$
- compute mutual information: $p(X, Y) \rightarrow I(X, Y)$

$$p(X=1, Y=1) = \frac{\text{\# of docs with topic } sports \text{ and containing score}}{\text{total \# of docs in the corpus}}$$

$$p(X=1, Y=0) = \frac{\text{\# of docs with topic } sports \text{ but not containing score}}{\text{total \# of docs in the corpus}}$$

$$\begin{aligned} I(X, Y) &= \sum_{x \in \{0,1\}} \sum_{y \in \{0,1\}} p(x,y) \log_2 \frac{p(x,y)}{p(x)p(y)} \\ &= 0.126 \end{aligned}$$

	$y=0$	$y=1$	$p(x)$
$x=0$	0.80	0.02	0.82
$x=1$	0.11	0.07	0.18
$p(y)$	0.91	0.09	

- for another pair, e.g. *what* vs. *sports*, $I(X, Y)=0.00007$
- *sports* is a keyword for the topic *sports*, *what* is not
- repeat the above procedure for all pairs of words and topics to identify keywords

KL Divergence (I)

KL divergence is a criterion to measure the difference between two probability distributions that have the same support

$$\text{KL}\left(p(x) \parallel q(x)\right) \triangleq \mathbb{E}_{x \sim p(x)} \left[\log \left(\frac{p(x)}{q(x)} \right) \right] = \sum_x p(x) \log \left(\frac{p(x)}{q(x)} \right)$$

Theorem 1

The KL divergence is always non-negative:

$$\text{KL}\left(p(x) \parallel q(x)\right) \geq 0$$

And $\text{KL}(p(x) \parallel q(x)) = 0$ iff $p(x) = q(x)$ holds almost everywhere.

proved as a corollary from Jensen's inequality

KL Divergence (II)

- $\text{KL}(q(x) \parallel p(x))$ represents the amount of information lost when we replace one probability distribution $p(x)$ with another distribution $q(x)$
- when using a simple model $q(x)$ to approximate a complicated model $p(x)$, the best-fit $q^*(x)$ may be derived as:

$$q^*(x) = \arg \min_{q(x)} \text{KL}(q(x) \parallel p(x))$$

- KL divergence vs. mutual information
 - $I(X, Y) = \text{KL}(p(x, y) \parallel p(x)p(y))$
 - mutual information can be viewed as the information gain from the assumption of independence

General Formulation for Optimization

the general formulation for all optimization problems:

$$\mathbf{x}^* = \arg \min_{\mathbf{x}} f(\mathbf{x})$$

subject to

$$h_i(\mathbf{x}) = 0 \quad (i = 1, 2, \dots, m)$$

$$g_j(\mathbf{x}) \leq 0 \quad (j = 1, 2, \dots, n)$$

where m equality constraints and n inequality constraints jointly define a non-empty feasible set Ω for the free variables \mathbf{x} .

- two special cases:
 - linear programming
 - convex optimization

Optimality Conditions

- **optimality conditions:** the necessary and/or sufficient conditions for any x^* to be an optimal solution of an optimization problem
- optimality conditions may help to derive the analytic solution to some simple optimization problems
- three cases from easy to hard:
 - unconstrained optimization
 - optimization under equality constraints
 - general optimization subject to both inequality and equality constraints

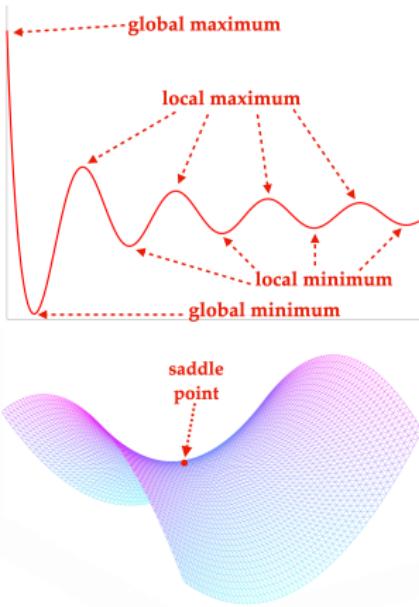
Optimality Conditions: Unconstrained Optimization (I)

$$\mathbf{x}^* = \arg \min_{\mathbf{x} \in \mathbb{R}^n} f(\mathbf{x})$$

- global minimum (maximum)
- local minimum (maximum)
- stationary point:

$$\nabla f(\hat{\mathbf{x}}) \triangleq \frac{\partial f(\mathbf{x})}{\partial \mathbf{x}} \Big|_{\mathbf{x}=\hat{\mathbf{x}}} = 0$$

- critical point: a stationary point or an undifferentiable point
- saddle point: a stationary point but not a local minimum (maximum)



Optimality Conditions: Unconstrained Optimization (II)

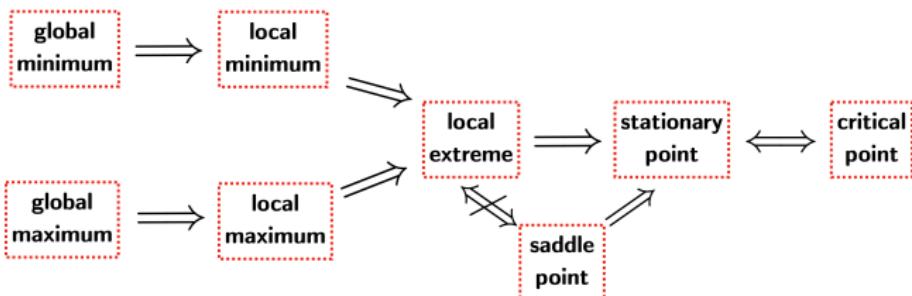


Figure: a diagram for any differentiable function

Theorem 2 (necessary condition for unconstrained optimization)

Assume $f(\mathbf{x})$ is differentiable everywhere. If \mathbf{x}^* is a local minimum, then \mathbf{x}^* must be a stationary point, i.e. the gradient vanishes at \mathbf{x}^* as $\nabla f(\mathbf{x}^*) = 0$.

Optimality Conditions: Unconstrained Optimization (III)

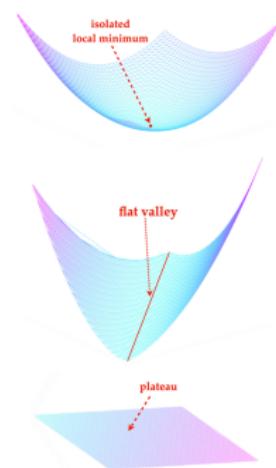
Hessian matrix: $\mathbf{H}(\mathbf{x}) = \left[\frac{\partial^2 f(\mathbf{x})}{\partial x_i \partial x_j} \right]_{n \times n}$

Theorem 3 (second order necessary condition)

Assume $f(\mathbf{x})$ is twice differentiable. If \mathbf{x}^* is a local minimum, then $\nabla f(\mathbf{x}^*) = 0$ and $\mathbf{H}(\mathbf{x}^*) \succeq 0$.

Theorem 4 (second order sufficient condition)

Assume $f(\mathbf{x})$ is twice differentiable. If a point \mathbf{x}^* satisfies $\nabla f(\mathbf{x}^*) = 0$ and $\mathbf{H}(\mathbf{x}^*) \succ 0$, then \mathbf{x}^* is an isolated local minimum.



- i) isolated minimum: $\mathbf{H}(\mathbf{x}) \succ 0$
- ii) flat valley: $\mathbf{H}(\mathbf{x}) \succeq 0$ and $\mathbf{H}(\mathbf{x}) \neq 0$
- iii) plateau: $\mathbf{H}(\mathbf{x}) = 0$

Optimality Conditions: Equality Constraints (I)

$$\mathbf{x}^* = \arg \min_{\mathbf{x}} f(\mathbf{x})$$

subject to $h_i(\mathbf{x}) = 0 \quad (i = 1, 2, \dots, m)$

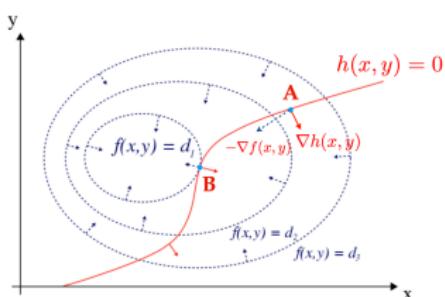
Theorem 5 (Lagrange necessary conditions)

Given $f(\mathbf{x})$ and $\{h_i(\mathbf{x})\}$ are differentiable.
If a point \mathbf{x}^* is a local optimum, then the gradients of these functions are linearly dependent at \mathbf{x}^* :

$$\nabla f(\mathbf{x}^*) + \sum_{i=1}^m \lambda_i \nabla h_i(\mathbf{x}^*) = 0$$

where $\lambda_i \in \mathbb{R}$ are the Lagrange multipliers.

stationary points of $f(\mathbf{x})$ may not satisfy the constraints.



$$\min_{x,y} f(x,y)$$

subject to

$$h(x,y) = 0$$

Optimality Conditions: Equality Constraints (II)

- Theorem 5 suggests **the method of Lagrange multipliers**
- introduce a Lagrange multiplier $\lambda_i \in \mathbb{R}$ for each equality constraint
- construct the *Lagrangian* function:

$$L(\mathbf{x}, \{\lambda_i\}) = f(\mathbf{x}) + \sum_{i=1}^m \lambda_i h_i(\mathbf{x})$$

- minimize the *Lagrangian* function w.r.t. \mathbf{x} and all Lagrange multipliers $\{\lambda_i\}$, converting a constrained optimization problem into an unconstrained one

Optimality Conditions: Inequality Constraints

$$\mathbf{x}^* = \arg \min_{\mathbf{x}} f(\mathbf{x})$$

subject to

$$h_i(\mathbf{x}) = 0 \quad (i = 1, 2, \dots, m)$$

$$g_j(\mathbf{x}) \leq 0 \quad (j = 1, 2, \dots, n)$$

- introduce a multiplier $\lambda_i \in \mathbb{R}$ for each equality constraint
- introduce a multiplier $\nu_i \geq 0$ for each inequality constraint
- construct a Lagrangian function:

$$L(\mathbf{x}, \{\lambda_i, \nu_j\}) = f(\mathbf{x}) + \sum_{i=1}^m \lambda_i h_i(\mathbf{x}) + \sum_{j=1}^n \nu_j g_j(\mathbf{x})$$

- optimize $L(\mathbf{x}, \{\lambda_i, \nu_j\})$ for the optimality conditions

Dual Problem and Strong Duality

- **Lagrange dual function:** defined as the lower-bound w.r.t. \mathbf{x}

$$L^*\left(\{\lambda_i, \nu_j\}\right) = \inf_{\mathbf{x} \in \Omega} L\left(\mathbf{x}, \{\lambda_i, \nu_j\}\right)$$

- generally $L^*\left(\{\lambda_i, \nu_j\}\right) \leq L\left(\mathbf{x}, \{\lambda_i, \nu_j\}\right) \leq f(\mathbf{x})$ for all $\mathbf{x} \in \Omega$
- **Lagrange dual problem:**

$$\{\lambda_i^*, \nu_j^*\} = \arg \max_{\{\lambda_i, \nu_j\}} L^*\left(\{\lambda_i, \nu_j\}\right)$$

subject to $\nu_j \geq 0$ for all $j = 1, 2, \dots, n$

- **strong duality** occurs if $L^*\left(\{\lambda_i^*, \nu_j^*\}\right) = f(\mathbf{x}^*)$, under which the dual problem is equivalent to the original problem

Optimality Conditions: KKT Conditions

Theorem 6 (KKT necessary conditions)

If \mathbf{x}^* and $\{\lambda_i^*, \nu_j^*\}$ is a saddle point of $L(\mathbf{x}, \{\lambda_i, \nu_j\})$, then \mathbf{x}^* is a local minimum. The saddle point satisfies the following conditions:

1 stationarity:

$$\nabla f(\mathbf{x}^*) + \sum_{i=1}^m \lambda_i^* \nabla h_i(\mathbf{x}^*) + \sum_{j=1}^n \nu_j^* \nabla g_j(\mathbf{x}^*) = 0$$

2 primal feasibility:

$$h_i(\mathbf{x}^*) = 0 \quad (i = 1, 2, \dots, m)$$

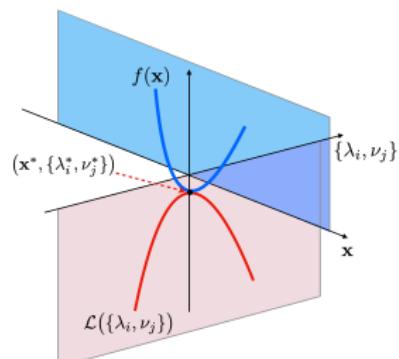
$$g_j(\mathbf{x}^*) \leq 0 \quad (j = 1, 2, \dots, n)$$

3 dual feasibility: $\nu_j^* \geq 0 \quad (j = 1, 2, \dots, n)$

4 complementary slackness:

$$\nu_j^* g_j(\mathbf{x}^*) = 0 \quad (j = 1, 2, \dots, n)$$

strong duality $\implies \mathbf{x}^*$ and $\{\lambda_i^*, \nu_j^*\}$ is a saddle point



Numerical Optimization Methods

- optimality conditions do not always yield a useful closed-form solution
- many optimization problems in machine learning require iterative numerical methods
- take the unconstrained optimization problem as example

$$\arg \min_{\mathbf{x} \in \mathbb{R}^n} f(\mathbf{x})$$

- numerical optimization methods
 - **zero-order methods**: using the function values alone, such as *grid search*
 - **first-order methods**: using the function values and gradients, such as *gradient descent method*
 - **second-order methods**: using the function values, gradients and Hessians, such as *Newton method*

First-order Methods: Gradient Descent Method

- the gradient points to a direction of the fastest increase of the function value
- gradient descent*: repeatedly move a small step along the direction of the negative gradient until converged

Gradient Descent Method

randomly choose $\mathbf{x}^{(0)}$, and set η_0

set $n = 0$

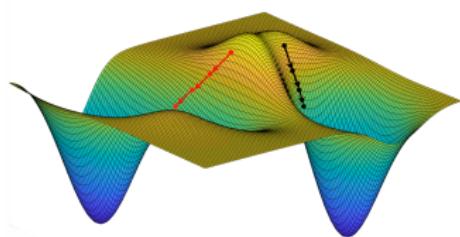
while not converged **do**

update: $\mathbf{x}^{(n+1)} = \mathbf{x}^{(n)} - \eta_n \nabla f(\mathbf{x}^{(n)})$

adjust: $\eta_n \rightarrow \eta_{n+1}$

$n = n + 1$

end while



First-order Methods: Stochastic Gradient Descent (I)

- the objective function $f(\mathbf{x})$ in machine learning can often be decomposed as a sum of homogeneous components:

$$f(\mathbf{x}) = \frac{1}{N} \sum_{i=1}^N f_i(\mathbf{x})$$

- e.g. $f_i(\mathbf{x})$ indicates the loss measure on each training sample
- when N is large, it is too expensive to compute

$$\nabla f(\mathbf{x}) = \frac{1}{N} \sum_{i=1}^N \nabla f_i(\mathbf{x})$$

- stochastic gradient descent:** estimate $\nabla f(\mathbf{x})$ using a random sample or a small subset (mini-batch) of random samples at each time

First-order Methods: Stochastic Gradient Descent (II)

Stochastic Gradient Descent (SGD)

randomly choose $\mathbf{x}^{(0)}$, and set η_0

set $n = 0$

while not converged **do**

 randomly choose a sample k

 update: $\mathbf{x}^{(n+1)} = \mathbf{x}^{(n)} - \eta_n \nabla f_k(\mathbf{x}^{(n)})$

 adjust: $\eta_n \rightarrow \eta_{n+1}$

$n = n + 1$

end while

First-order Methods: Stochastic Gradient Descent (III)

Mini-batch SGD

randomly choose $\mathbf{x}^{(0)}$, and set η_0

set $n = 0$

while not converged **do**

 randomly shuffle all training samples into mini-batches

for each mini-batch B **do**

 update: $\mathbf{x}^{(n+1)} = \mathbf{x}^{(n)} - \frac{\eta_n}{|B|} \sum_{k \in B} \nabla f_k(\mathbf{x})$

 adjust $\eta_n \rightarrow \eta_{n+1}$

$n = n + 1$

end for

end while

Second-order Methods: Newton method

- expand $f(\mathbf{x})$ at any fixed \mathbf{x}_0 according to the Taylor's theorem

$$f(\mathbf{x}) \approx f(\mathbf{x}_0) + (\mathbf{x} - \mathbf{x}_0)^T \nabla f(\mathbf{x}_0) + \frac{1}{2} (\mathbf{x} - \mathbf{x}_0)^T \mathbf{H}(\mathbf{x}_0) (\mathbf{x} - \mathbf{x}_0)$$

- $\nabla f(\mathbf{x}) = \frac{\partial f(\mathbf{x})}{\partial \mathbf{x}} = 0 \implies \mathbf{x}^* = \mathbf{x}_0 - \mathbf{H}^{-1}(\mathbf{x}_0) \nabla f(\mathbf{x}_0)$
- **Newton method** uses the updating rule at each iteration:

$$\mathbf{x}^{(n+1)} = \mathbf{x}^{(n)} - \mathbf{H}^{-1}(\mathbf{x}^{(n)}) \nabla f(\mathbf{x}^{(n)})$$

- Newton method is not feasible in machine learning since it is too costly to invert a Hessian matrix
- **quasi-Newton methods** aim to approximate the Hessian matrix, e.g. DFP, BFGS, Quickprop, Hessian-free, etc.