

# Chapter 8

## Neural Networks

supplementary slides to  
*Machine Learning Fundamentals*  
©Hui Jiang 2020  
published by Cambridge University Press

August 2020



CAMBRIDGE  
UNIVERSITY PRESS

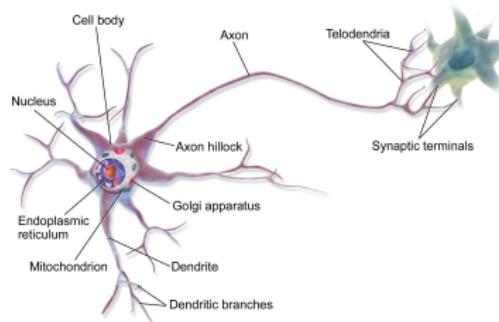
# Outline

- 1 Artificial Neural Networks**
- 2 Neural Network Structures**
- 3 Learning Algorithms for Neural Networks**
- 4 Heuristics and Tricks for Optimization**
- 5 End-to-End Learning**

# Biological Neuronal Networks



(a) neuronal networks

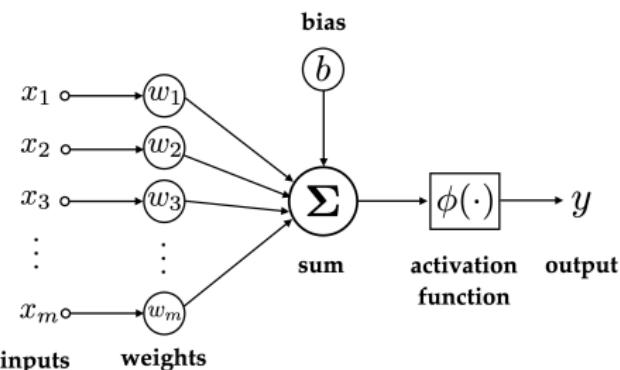


(b) biological neuron

- brain: a large number of inter-connected neurons
  - neuron: axon, dendrites and synapse
  - mechanisms of biological neuronal networks

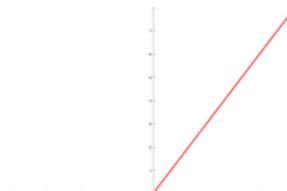
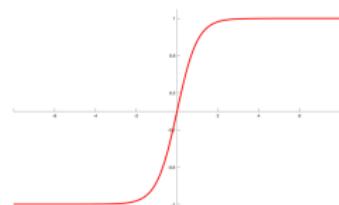
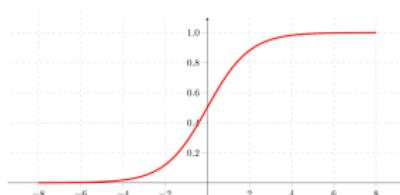
## Artificial Neural Networks (ANNs)

- motivated by biological neuronal networks
  - artificial neuron: a simplified computational model to simulate a biological neuron  $y = \phi(\sum_i w_i x_i + b)$ 
    - nonlinear activation function: sigmoid, tanh, ReLU, etc.



- ANNs consist of a large number of artificial neurons

# Nonlinear Activation Functions



$$\text{sigmoid} : y = \frac{1}{1 + e^{-x}}$$

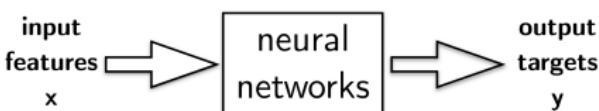
$$\tanh : y = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

$$\text{ReLU} : y = \max(0, x)$$

- sigmoid:  $(0, 1)$ , monotonically increasing, differentiable everywhere
  - tanh:  $(-1, 1)$ , monotonically increasing, differentiable everywhere
  - ReLU:  $[0, \infty)$ , monotonically non-decreasing, unbounded

# Neural Networks: Mathematical Justification

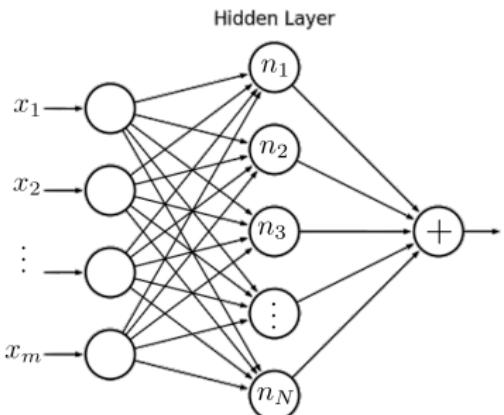
- neural networks are primarily used as a function approximator



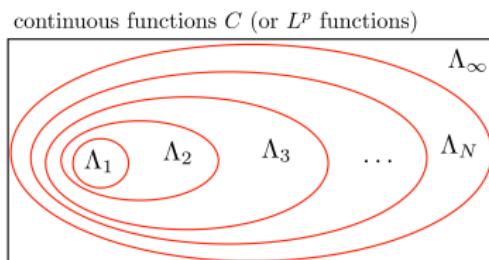
- what is the modeling power of neural networks?
  - linear functions vs. nonlinear functions
  - $f(\mathbf{x})$  is an  $L^p$  function ( $\forall p > 0$ ) iff  $\int_{\mathbf{x}} |f(\mathbf{x})|^p d\mathbf{x} < \infty$
  - including either energy-limited functions, or bounded functions on finite-domain
  - e.g. all  $L^2$  functions ( $p = 2$ ) form a Hilbert space, consisting of all functions arising from any physical process

# Neural Networks: Universal Approximator (I)

multilayer perceptrons (MLP): a simple structure for neural nets, containing only one hidden layer between input and output



(c) MLP



$$\Lambda_1 \subset \Lambda_2 \subset \Lambda_3 \cdots \subset \Lambda_N \subset \cdots \subset \Lambda_\infty \equiv C(\text{or } L^p)$$



# Neural Networks: Universal Approximator (II)

- MLPs are universal function approximators

## Theorem 1

Denote all continuous functions on  $\mathbb{R}^m$  as  $C$ . If the nonlinear activation function  $\phi(\cdot)$  is continuous, bounded and non-constant, then  $\Lambda_N$  is dense in  $C$  as  $N \rightarrow \infty$ , i.e.  $\lim_{N \rightarrow \infty} \Lambda_N = C$ .

## Theorem 2

Denote all  $L^p$  functions on  $\mathbb{R}^m$  as  $L^p$ . If the ReLU function is used as the activation function  $\phi(\cdot)$ , then  $\Lambda_N$  is dense in  $L^p$  as  $N \rightarrow \infty$ , i.e.  $\lim_{N \rightarrow \infty} \Lambda_N = L^p$ .

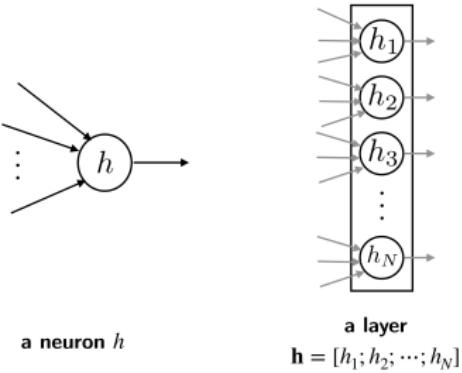
- applicable to many other neural network structures

# Neural Network Structures

- Neurons vs. Layers of Neurons
- Building Blocks
  - full connection, convolution
  - nonlinear activation, softmax, max-pooling
  - normalization
  - time-delayed feedback
  - tapped delay line
  - attention
- Case Studies:
  - 1 fully-connected deep neural networks (DNNs)
  - 2 convolutional neural networks (CNNs)
  - 3 recurrent neural networks (RNNs)
  - 4 transformers

# Neurons vs. Layers of Neurons

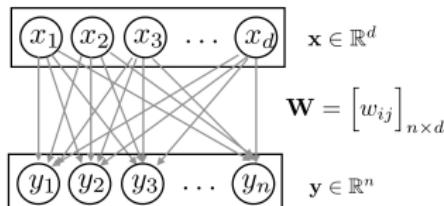
- a neuron: mathematically represents a variable in computation
- convenient to group relevant neurons as a layer
- a layer of neurons: represents a vector in computation
- neural nets are constructed by arranging layers of neurons



# Building Blocks to Connect Layers (I)

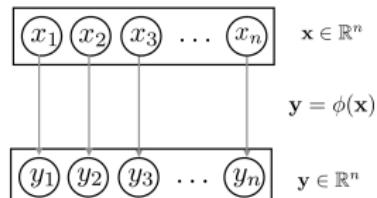
## ■ full connection: $y = \mathbf{W}\mathbf{x} + \mathbf{b}$

- $\mathbf{W} \in \mathbb{R}^{n \times d}$  and  $\mathbf{b} \in \mathbb{R}^n$  denote all parameters in a full connection
- $n \times (d + 1)$  parameters
- computational complexity is  $O(n \times d)$
- mainly used for universal function approximation



## ■ nonlinear activation: $y = \phi(\mathbf{x})$

- $\phi(\cdot)$ : ReLU, sigmoid or tanh
- no learnable parameter in this connection
- used to introduce nonlinearity



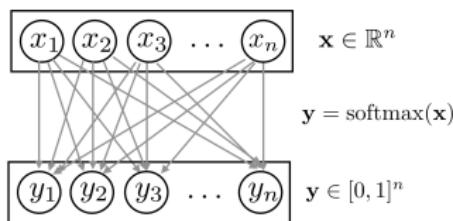
## Building Blocks to Connect Layers (II)

### ■ softmax

$$\mathbf{y} = \text{softmax}(\mathbf{x})$$

where  $y_i = \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}}$  for all  $i$

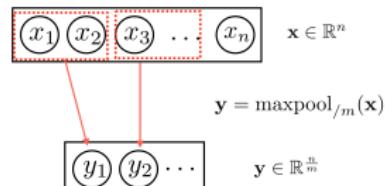
- o no learnable parameter in this connection
  - o used to generate probability-like outputs



## ■ max-pooling

$$\mathbf{y} = \text{maxpool}_{/m}(\mathbf{x}) \quad (\mathbf{x} \in \mathbb{R}^n, \mathbf{y} \in \mathbb{R}^{\frac{n}{m}})$$

- no learnable parameter in this connection
  - used to reduce the layer size
  - make the output less sensitive to small translation variations



# Building Blocks to Connect Layers (III)

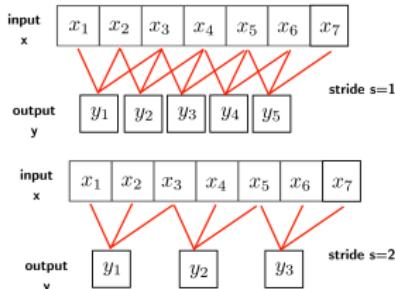
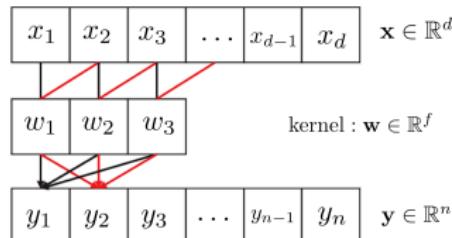
## ■ convolution:

$$\mathbf{y} = \mathbf{x} * \mathbf{w} \quad (\mathbf{x} \in \mathbb{R}^d, \mathbf{w} \in \mathbb{R}^f, \mathbf{y} \in \mathbb{R}^n)$$

where  $y_j = \sum_{i=1}^f w_i \times x_{j+i-1} \quad (\forall j)$

- kernel  $\mathbf{w}$  represents  $f$  learnable parameters
- computational complexity:  $O(d \times f)$
- output neurons are  $n = d - f + 1$  but can be adjusted by zero-padding and striding
- convolution vs. full connection

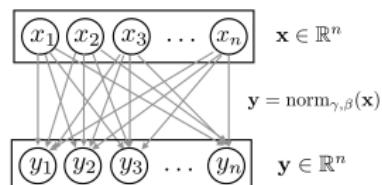
- 1 locality modelling: only capture a local feature
- 2 weight sharing:  $f(< d)$  weights (vs.  $d \times n$  weights in full connection)



# Building Blocks to Connect Layers (IV)

## ■ normalization

- normalize the dynamic ranges of neurons
- smooth out the loss surface to facilitate optimization



1 *batch normalization*:  $\mathbf{y} = \text{BN}_{\gamma, \beta}(\mathbf{x})$

$$(1) \text{ normalize: } \hat{x}_i = \frac{x_i - \mu_B(i)}{\sqrt{\sigma_B^2(i) + \epsilon}} \quad (2) \text{ re-scaling: } y_i = \gamma_i \hat{x}_i + \beta_i$$

where  $\mu_B(i)$  and  $\sigma_B^2(i)$  denote the sample mean and the sample variance over the current mini-batch

2 *layer normalization*:  $\mathbf{y} = \text{LN}_{\gamma, \beta}(\mathbf{x})$

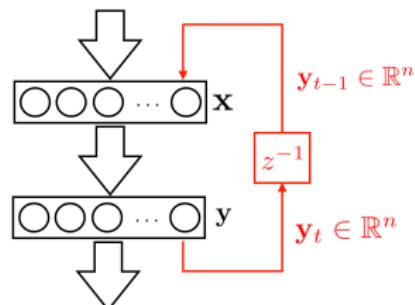
where local statistics are estimated over all dimensions in each input vector  $\mathbf{x}$

# Building Blocks to Connect Layers (V)

## ■ time-delayed feedback

$$\mathbf{y}_{t-1} = z^{-1}(\mathbf{y}_t)$$

- $z^{-1}$  indicates a time-delay unit, which is physically implemented as a memory unit
- recurrent neural networks (RNNs) use feedback to memorize the history
- feedback paths introduce circles in nets



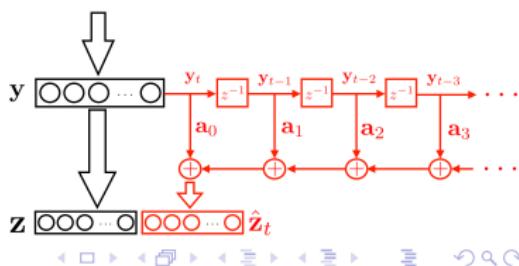
## ■ tapped delay line

- stored in a line of memory units
- linearly combined to feed forward

$$\hat{\mathbf{z}}_t = \sum_{i=0}^{L-1} \mathbf{a}_i \otimes \mathbf{y}_{t-i}$$

where  $\{\mathbf{a}_i\}$  are learnable parameters

- no feedback path  $\Rightarrow$  non-recurrent structures to memorize the history



# Building Blocks to Connect Layers (V): Attention (1)

- **attention:** use time-variant scalar coefficients in tapped delay lines

- 1 tapped-delay-line is long enough to store entire sequence
- 2 introduce an attention function  $g()$

$$g(\mathbf{q}_t, \mathbf{k}_t) \triangleq [c_0(t) \ c_1(t) \ \cdots \ c_{L-1}(t)]^T$$

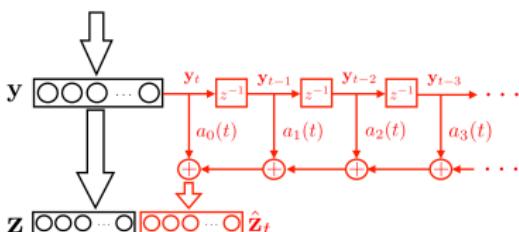
- o  $\mathbf{q}_t \in \mathbb{R}^l$ : query vector at time  $t$
- o  $\mathbf{k}_t \in \mathbb{R}^l$ : key vector at time  $t$

- 3 normalize to one by softmax

$$\mathbf{a}_t = \text{softmax}(g(\mathbf{q}_t, \mathbf{k}_t))$$

- 4 linearly combined at each time  $t$

$$\hat{\mathbf{z}}_t = \sum_{i=0}^{L-1} a_i(t) \mathbf{y}_{t-i} = [\mathbf{y}_t \ \mathbf{y}_{t-1} \ \cdots \ \mathbf{y}_{t-L+1}] \ \mathbf{a}_t$$



## Building Blocks to Connect Layers (VI): Attention (2)

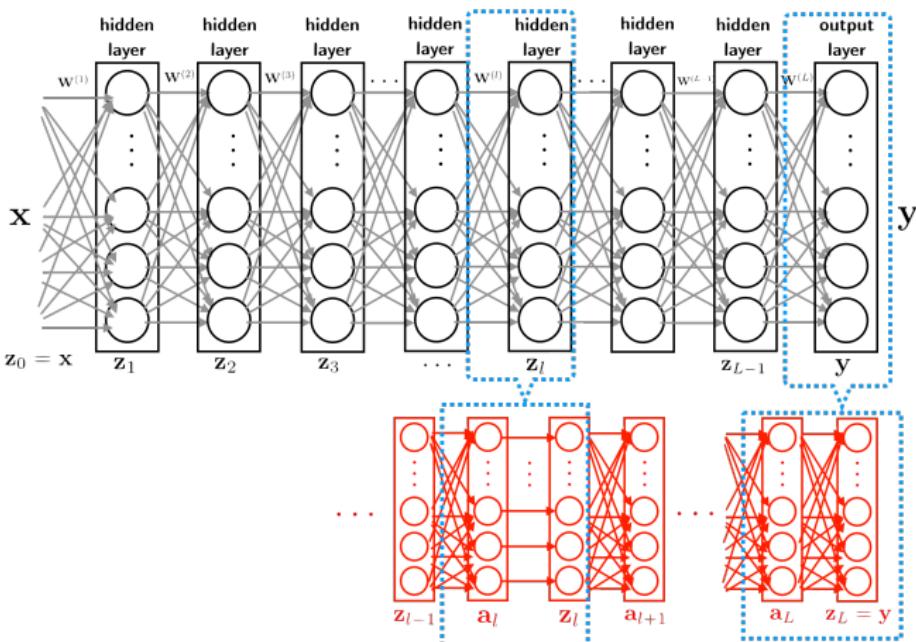
- use a matrix form to represent attention for all time instances
- value matrix:  $\mathbf{V} = [\mathbf{y}_T \ \mathbf{y}_{T-1} \ \cdots \ \mathbf{y}_1]_{n \times T}$
- query matrix:  $\mathbf{Q} \triangleq [\mathbf{q}_T \ \mathbf{q}_{T-1} \ \cdots \ \mathbf{q}_1]_{l \times T}$
- key matrix:  $\mathbf{K} \triangleq [\mathbf{k}_T \ \mathbf{k}_{T-1} \ \cdots \ \mathbf{k}_1]_{l \times T}$
- attention in a compact form:

$$\hat{\mathbf{Z}} = \mathbf{V} \text{ softmax}\left(g(\mathbf{Q}, \mathbf{K})\right)$$

where softmax is applied to  $g(\mathbf{Q}, \mathbf{K}) \in \mathbb{R}^{T \times T}$  column-wise

- attention represents a very flexible and complex computation in neural networks, depending on how to choose the four elements:  $\mathbf{V}$ ,  $\mathbf{Q}$ ,  $\mathbf{K}$  and  $g(\cdot)$

# Case Study (I): Fully-Connected Deep Neural Networks (1)



# Case Study (I): Fully-Connected Deep Neural Networks (2)

## Forward Pass of a fully-connected DNN

- 1 For the input layer:  $\mathbf{z}_0 = \mathbf{x}$
- 2 For each hidden layer  $l = 1, 2, \dots, L - 1$ :

$$\mathbf{a}_l = \mathbf{W}^{(l)} \mathbf{z}_{l-1} + \mathbf{b}^{(l)}$$

$$\mathbf{z}_l = \text{ReLU}(\mathbf{a}_l)$$

- 3 For the output layer:

$$\mathbf{a}_L = \mathbf{W}^{(L)} \mathbf{z}_{L-1} + \mathbf{b}^{(L)}$$

$$\mathbf{y} = \mathbf{z}_L = \text{softmax}(\mathbf{a}_L)$$

## Case Study (II): Convolutional Neural Networks

- convolutional neural networks (CNNs) are currently the dominant model for images/videos
- CNNs mainly rely on the basic convolution sum
- extension #1: allow multiple feature plies in input
- extension #2: allow multiple kernels
- extension #3: allow multiple input dimensions
- extension #4: stack many convolution layers
- typical CNN architectures:
  - AlexNet, VGG, ResNet, etc.

# From Convolution Sum to CNNs (1)

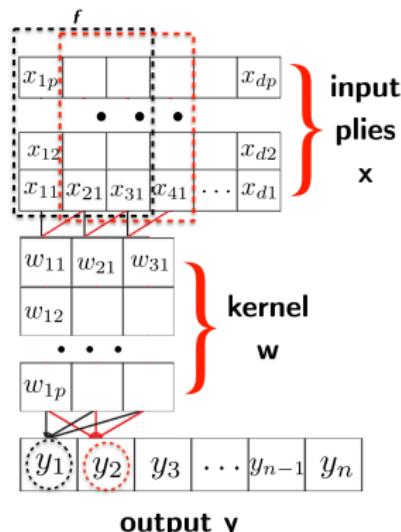
**extension #1:** allow multiple feature plies in input  $\mathbf{x}$

- each input position contains  $p$  feature plies (e.g. R/G/B in color images)
- extend kernel to  $p$  plies ( $p \times f$  weights)

$$y_j = \sum_{k=1}^p \sum_{i=1}^f w_{i,k} \times x_{j+i-1,k} \quad (\forall j = 1, 2, \dots, n)$$

$$\mathbf{y} = \mathbf{x} * \mathbf{w} \quad (\mathbf{x} \in \mathbb{R}^{p \times d}, \mathbf{w} \in \mathbb{R}^{p \times f}, \mathbf{y} \in \mathbb{R}^n)$$

- computational complexity:  $O(d \cdot f \cdot p)$
- zero-padding and striding
- locality modeling, weight sharing



# From Convolution Sum to CNNs (2)

**extension #2:** allow multiple kernels for more local features

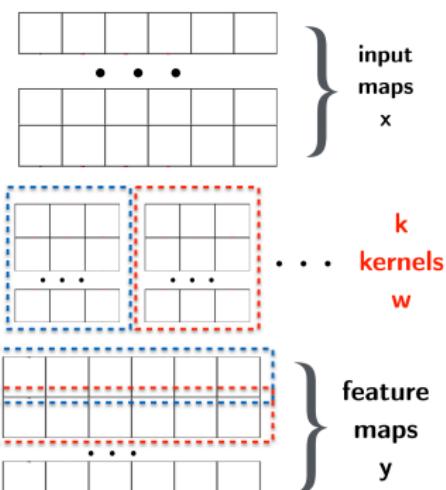
- a kernel captures only one local feature
- extend to  $k$  kernels ( $p \times f \times k$  weights)
- output is a  $k \times n$  feature map

$$y_{j_1, j_2} = \sum_{i_2=1}^p \sum_{i_1=1}^f w_{i_1, i_2, j_2} \times x_{j_1+i_1-1, i_2}$$

$$(\forall j_1 = 1, \dots, n; j_2 = 1, \dots, k)$$

$$\mathbf{y} = \mathbf{x} * \mathbf{w} \quad (\mathbf{x} \in \mathbb{R}^{p \times d}, \mathbf{w} \in \mathbb{R}^{p \times f \times k}, \mathbf{y} \in \mathbb{R}^{k \times n})$$

- computational complexity:  $O(d \cdot f \cdot p \cdot k)$
- zero-padding and striding
- locality modeling, weight sharing



# From Convolution Sum to CNNs (3)

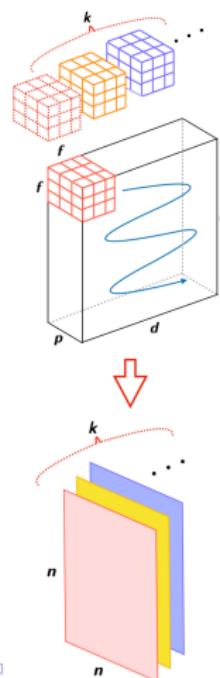
- **extension #3:** allow multiple input dimensions
- expand input dimension to handle multi-dim data, e.g. images (2D) and videos (3D)
- for 2D images, each input  $\mathbf{x}$  is a  $d \times d \times p$  tensor, extend each kernel into an  $f \times f \times p$  tensor, output is an  $n \times n \times k$  feature map

$$y_{j_1, j_2, j_3} = \sum_{i_3=1}^p \sum_{i_2=1}^f \sum_{i_1=1}^f w_{i_1, i_2, i_3, j_3} \times x_{j_1+i_1-1, j_2+i_2-1, i_3}$$

$$(j_1 = 1, \dots, n; j_2 = 1, \dots, n; j_3 = 1, \dots, k)$$

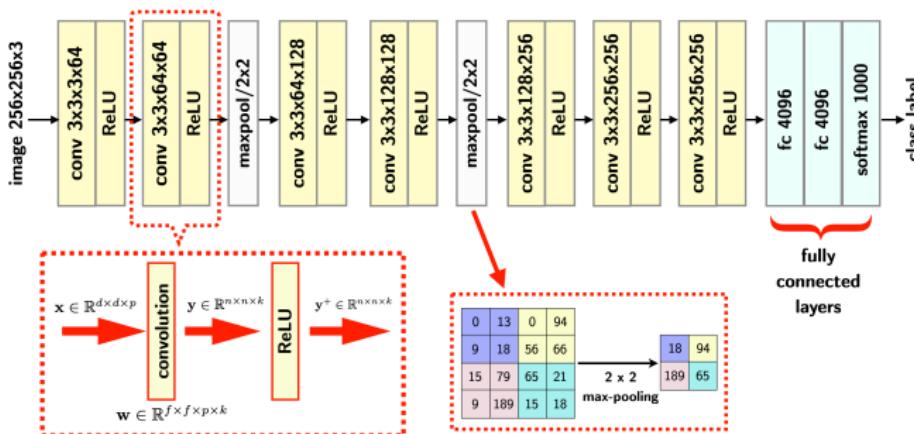
$$\mathbf{y} = \mathbf{x} * \mathbf{w} \quad (\mathbf{x} \in \mathbb{R}^{d \times d \times p}, \mathbf{w} \in \mathbb{R}^{f \times f \times p \times k}, \mathbf{y} \in \mathbb{R}^{n \times n \times k})$$

- computational complexity:  $O(d^2 \cdot f^2 \cdot p \cdot k)$
- locality modeling: capture 2D local features



# From Convolution Sum to CNNs (4)

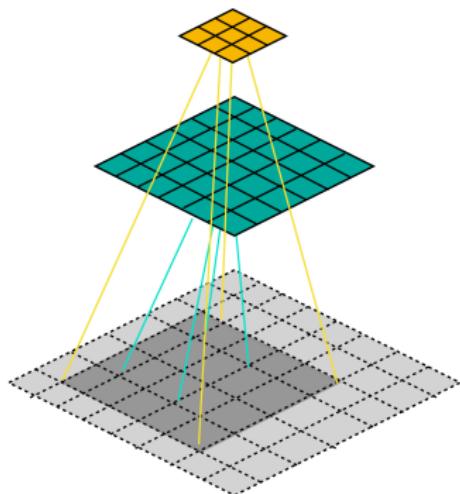
**extension #4:** stack many convolution layers to form CNNs



- stacked convolution layers: hierarchical visual feature extraction
- fully-connected layers: a universal function approximator to map these features to the target labels

# Convolutional Neural Networks (CNNs)

- locality modelling  $\implies$  hierarchical modeling
  - recursively combine local features
  - **receptive fields** in CNN: broaden in upper layers
- CNNs are dominant in image classification, segmentation, generation
- typical CNN architectures:
  - AlexNet, VGG, ResNet, etc.
  - *ResNet*: a very deep structure with shortcut paths



## Case Study (III): Recurrent Neural Network (RNN)

- use a simple RNN to process a sequence of input vectors:  $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T\}$
- for all  $t = 1, 2, \dots, T$

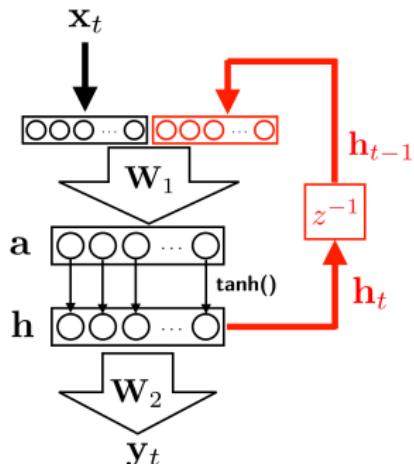
$$\mathbf{a}_t = \mathbf{W}_1[\mathbf{x}_t; \mathbf{h}_{t-1}] + \mathbf{b}_1$$

$$\mathbf{h}_t = \tanh(\mathbf{a}_t)$$

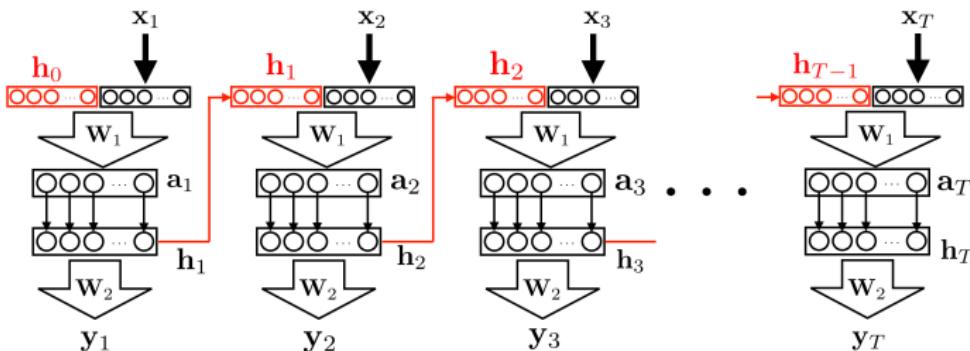
$$\mathbf{y}_t = \mathbf{W}_2 \mathbf{h}_t + \mathbf{b}_2$$

where  $\mathbf{W}_1$ ,  $\mathbf{b}_1$ ,  $\mathbf{W}_2$  and  $\mathbf{b}_2$  are all RNN parameters

- RNN generates an output sequence:  $\{\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_T\}$



# Case Study (III): Recurrent Neural Network (RNN)



- an RNN can be unfolded into a non-recurrent structure
- RNNs fail to capture long-term dependency due to long traversal paths in the deep structures
- more effective RNN structures, e.g. LSTMs, GRUs, HORNNS

## Case Study (IV): Transformer (1)

use a particular attention mechanism to directly map an input sequence to an output sequence

$$\mathbf{X} = [\mathbf{x}_T \ \cdots \ \mathbf{x}_2 \ \mathbf{x}_1] \ \longmapsto \ \mathbf{Z} = [\mathbf{z}_T \ \cdots \ \mathbf{z}_2 \ \mathbf{z}_1]$$

- 1 choose query matrix  $\mathbf{Q}$ , key matrix  $\mathbf{K}$  and value matrix  $\mathbf{V}$  as:

$$\mathbf{Q} = \mathbf{A}\mathbf{X} \quad \mathbf{K} = \mathbf{B}\mathbf{X} \quad \mathbf{V} = \mathbf{C}\mathbf{X}$$

where  $\mathbf{A}, \mathbf{B} \in \mathbb{R}^{l \times d}$ ;  $\mathbf{C} \in \mathbb{R}^{o \times d}$ ;  $\mathbf{Q}, \mathbf{K} \in \mathbb{R}^{l \times T}$  and  $\mathbf{V} \in \mathbb{R}^{o \times T}$

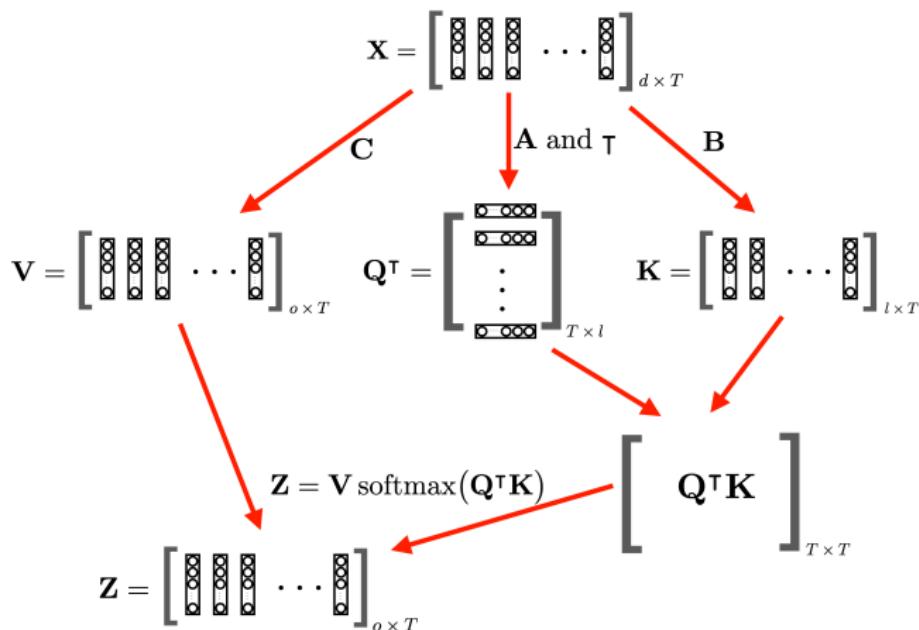
- 2 define the attention function as a bilinear function:

$$g(\mathbf{Q}, \mathbf{K}) = \mathbf{Q}^\top \mathbf{K} \quad (\in \mathbb{R}^{T \times T})$$

- 3 transformer as attention:

$$\mathbf{Z} = (\mathbf{C}\mathbf{X}) \text{ softmax}\left( (\mathbf{A}\mathbf{X})^\top (\mathbf{B}\mathbf{X}) \right)$$

## Case Study (IV): Transformer (2)



two enhancements:

- 1 use multiple heads in each transformer
- 2 stack more transformer layers to form a deep structure

# Case Study (IV): Transformer (3)

## Multi-head Transformer

Choose  $d = 512$ ,  $o = 64$ , a multi-head transformer will transform an input sequence  $\mathbf{X} \in \mathbb{R}^{512 \times T}$  into  $\mathbf{Y} \in \mathbb{R}^{n \times T}$ :

- multi-head transformer: use 8 sets of parameters  $\mathbf{A}^{(j)}, \mathbf{B}^{(j)} \in \mathbb{R}^{l \times 512}, \mathbf{C}^{(j)} \in \mathbb{R}^{64 \times 512}$  ( $j = 1, 2, \dots, 8$ )
- for  $j = 1, 2, \dots, 8$ :

$$\mathbf{Z}^{(j)} \in \mathbb{R}^{64 \times T} = (\mathbf{C}^{(j)} \mathbf{X}) \text{ softmax}\left((\mathbf{A}^{(j)} \mathbf{X})^\top (\mathbf{B}^{(j)} \mathbf{X})\right)$$

- concatenate all heads:  $\mathbf{Z} \in \mathbb{R}^{512 \times T} = \text{concat}(\mathbf{Z}^{(1)}, \mathbf{Z}^{(2)}, \dots, \mathbf{Z}^{(8)})$
- apply nonlinearity:  $\mathbf{Y} = \text{feedforward}\left(\text{LN}_{\gamma, \beta}(\mathbf{X} + \mathbf{Z})\right)$

# Learning Neural Networks

- Loss Function
- Optimization Method: SGD
- Automatic Differentiation
  - full connection
  - nonlinear activation
  - softmax
  - max-pooling
  - convolution
  - normalization
- Error Backpropagation Examples:
  - fully-connected deep neural networks

# Loss Function

- once network structure is determined, a neural network can be viewed as a multivariate and vector-valued function as:

$$\mathbf{y} = f(\mathbf{x}; \mathbb{W})$$

where  $\mathbb{W}$  to denote all network parameters

- learn  $\mathbb{W}$  from a training set of input-output pairs:

$$\mathcal{D}_N = \left\{ (\mathbf{x}_1, \mathbf{r}_1), (\mathbf{x}_2, \mathbf{r}_2), \dots, (\mathbf{x}_N, \mathbf{r}_N) \right\}$$

- mean square error (MSE) for regression problems

$$Q_{\text{MSE}}(\mathbb{W}; \mathcal{D}_N) = \sum_{i=1}^N \|f(\mathbf{x}_i; \mathbb{W}) - \mathbf{r}_i\|^2$$

- cross-entropy (CE) error for classification problems

$$Q_{\text{CE}}(\mathbb{W}; \mathcal{D}_N) = - \sum_{i=1}^N \ln [\mathbf{y}_i]_{r_i} = - \sum_{i=1}^N \ln [f(\mathbf{x}_i; \mathbb{W})]_{r_i}$$

# Optimization Method: mini-batch SGD

mini-batch SGD to learn neural networks

randomly initialize  $\mathbb{W}^{(0)}$ ; set  $\eta_0$ ,  $n = 0$  and  $t = 0$

**while** not converged **do**

    randomly shuffle training data into mini-batches

**for** each mini-batch  $B$  **do**

**for** each  $\mathbf{x} \in B$  **do**

            compute the gradient:  $\frac{\partial Q(\mathbb{W}^{(n)}; \mathbf{x})}{\partial \mathbb{W}}$

**end for**

        update model:  $\mathbb{W}^{(n+1)} = \mathbb{W}^{(n)} - \frac{\eta_t}{|B|} \sum_{\mathbf{x} \in B} \frac{\partial Q(\mathbb{W}^{(n)}; \mathbf{x})}{\partial \mathbb{W}}$

$n = n + 1$

**end for**

    adjust  $\eta_t \rightarrow \eta_{t+1}$

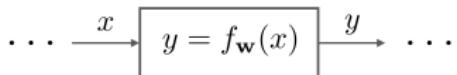
$t = t + 1$

**end while**

# Automatic Differentiation (I)

- how to efficiently compute gradients for arbitrary networks?
- automatic differentiation (AD), a.k.a. error back-propagation:
  - the most efficient for any network structure by systematically applying the chain rule

a simple example:



given any objective  
function  $Q(\cdot)$

- 1 define the error signal:  $e = \frac{\partial Q}{\partial y}$
- 2 derive the gradient by local computations:

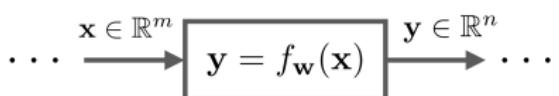
$$\frac{\partial Q}{\partial \mathbf{w}} = \frac{\partial Q}{\partial y} \frac{\partial y}{\partial \mathbf{w}} = e \frac{\partial f_{\mathbf{w}}(x)}{\partial \mathbf{w}}$$

- 3 back-propagate the error signal:

$$\frac{\partial Q}{\partial x} = \frac{\partial Q}{\partial y} \frac{\partial y}{\partial x} = e \frac{df_{\mathbf{w}}(x)}{dx}$$

# Automatic Differentiation (II)

extend AD to a vector-input and vector-output module:



$$\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} \quad \mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_m \end{bmatrix} \quad \mathbf{w} = \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_k \end{bmatrix}$$

compute two Jacobian matrices:

$$J_{\mathbf{w}} = \left[ \frac{\partial y_j}{\partial w_i} \right]_{k \times n}$$

$$J_{\mathbf{x}} = \left[ \frac{\partial y_j}{\partial x_i} \right]_{m \times n}$$

- given the error signal  $\mathbf{e} \triangleq \frac{\partial Q}{\partial \mathbf{y}}$  ( $\in \mathbb{R}^n$ )

## 1. local gradients:

$$\frac{\partial Q}{\partial \mathbf{w}} = J_{\mathbf{w}} \mathbf{e}$$

## 2. back-propagation:

$$\frac{\partial Q}{\partial \mathbf{x}} = J_{\mathbf{x}} \mathbf{e}$$

# Automatic Differentiation (III)

- **full connection** from  $\mathbf{x} \in \mathbb{R}^d$  to output  $\mathbf{y} \in \mathbb{R}^n$ :

$$\mathbf{y} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

where  $\mathbf{W} \in \mathbb{R}^{n \times d}$  and  $\mathbf{b} \in \mathbb{R}^n$

- back-propagation:

$$J_{\mathbf{x}} = \left[ \frac{\partial y_j}{\partial x_i} \right]_{d \times n} = \mathbf{W}^T \quad \Rightarrow \quad \frac{\partial Q}{\partial \mathbf{x}} = \mathbf{W}^T \mathbf{e}$$

- local gradients:

$$\frac{\partial Q}{\partial \mathbf{W}} = \begin{bmatrix} \frac{\partial Q}{\partial y_1} \\ \vdots \\ \frac{\partial Q}{\partial y_n} \end{bmatrix} \mathbf{x}^T = \mathbf{e} \mathbf{x}^T \quad \frac{\partial Q}{\partial \mathbf{b}} = \mathbf{e}$$

# Automatic Differentiation (III)

- **nonlinear activation** from  $\mathbf{x}$  ( $\in \mathbb{R}^n$ ) to  $\mathbf{y}$  ( $\in \mathbb{R}^n$ ):

$$\mathbf{y} = \phi(\mathbf{x})$$

- no learnable parameters  $\implies$  no local gradients
- back-propagation:

$$\frac{\partial Q}{\partial \mathbf{x}} = \mathbf{J}_{\mathbf{x}} \mathbf{e} = \phi'(\mathbf{x}) \odot \mathbf{e}$$

where  $\odot$  denotes element-wise multiplication

- for ReLU activation:  $\frac{\partial Q}{\partial \mathbf{x}} = H(\mathbf{x}) \odot \mathbf{e}$
- for sigmoid activation:  $\frac{\partial Q}{\partial \mathbf{x}} = l(\mathbf{x}) \odot (1 - l(\mathbf{x})) \odot \mathbf{e}$

# Automatic Differentiation (IV)

- **softmax:** mapping an  $n$ -dimensional vector  $\mathbf{x}$  ( $\in \mathbb{R}^n$ ) into another  $n$ -dimensional vector  $\mathbf{y}$  inside the hypercube  $[0, 1]^n$ , with  $y_j = \frac{e^{x_j}}{\sum_{i=1}^n e^{x_i}}$  for all  $i = 1, 2, \dots, n$
- no learnable parameters  $\implies$  no local gradients
- the Jacobian matrix

$$\mathbf{J}_{\mathbf{x}} = \left[ \frac{\partial y_j}{\partial x_i} \right]_{n \times n} = \begin{bmatrix} y_1(1 - y_1) & -y_1 y_2 & \cdots & -y_1 y_n \\ -y_1 y_2 & y_2(1 - y_2) & \cdots & -y_2 y_n \\ \vdots & \vdots & \ddots & \vdots \\ -y_1 y_n & -y_2 y_n & \cdots & y_n(1 - y_n) \end{bmatrix}_{n \times n}$$

- back-propagation:

$$\frac{\partial Q}{\partial \mathbf{x}} = \mathbf{J}_{\mathbf{x}} \mathbf{e}$$

# Automatic Differentiation (V): Convolution (1)

- **convolution:** mapping an input vector  $\mathbf{x} \in \mathbb{R}^d$  to an output vector  $\mathbf{y} \in \mathbb{R}^n$  by  $\mathbf{y} = \mathbf{x} * \mathbf{w}$  with  $\mathbf{w} \in \mathbb{R}^f$ , with

$$y_j = \sum_{i=1}^f w_i \times x_{j+i-1} \quad j = 1, 2 \dots, n$$

- the Jacobian matrix  $\mathbf{J}_{\mathbf{x}}$ :

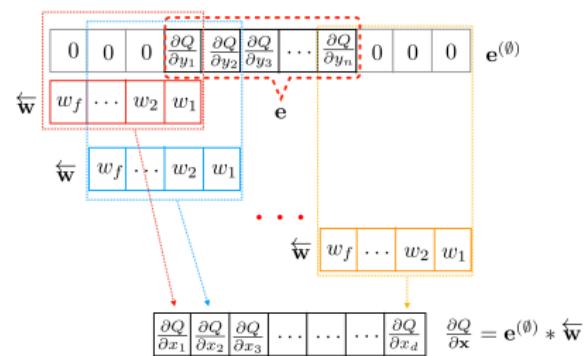
$$\mathbf{J}_{\mathbf{x}} = \left[ \frac{\partial y_j}{\partial x_i} \right]_{d \times n} = \begin{bmatrix} w_1 & & & \\ w_2 & w_1 & & \\ \vdots & \vdots & \ddots & \\ w_f & w_{f-1} & \ddots & w_1 \\ & w_f & & w_2 \\ & & \ddots & \vdots \\ & & & w_f \end{bmatrix}_{d \times n}$$

# Automatic Differentiation (V): Convolution (2)

- back-propagation by convolution:

$$\frac{\partial Q}{\partial \mathbf{x}} = \mathbf{J}_{\mathbf{x}} \mathbf{e} = \begin{bmatrix} w_1 \frac{\partial Q}{\partial y_1} \\ w_2 \frac{\partial Q}{\partial y_1} + w_1 \frac{\partial Q}{\partial y_2} \\ \vdots \\ w_f \frac{\partial Q}{\partial y_n} \end{bmatrix}$$

$$\triangleq \mathbf{e}^{(\emptyset)} * \overleftarrow{\mathbf{w}}$$



- computing local gradients by convolution:

$$\mathbf{J}_{\mathbf{w}} = \left[ \frac{\partial y_j}{\partial w_i} \right]_{f \times n} = \begin{bmatrix} x_1 & x_2 & \cdots & x_n \\ x_2 & x_3 & \cdots & x_{n+1} \\ \vdots & \vdots & \ddots & \vdots \\ x_f & x_{f+1} & \cdots & x_{n+f-1} \end{bmatrix}_{f \times n} \implies \frac{\partial Q}{\partial \mathbf{w}} = \mathbf{J}_{\mathbf{w}} \mathbf{e} \triangleq \mathbf{x} * \mathbf{e}$$

# Automatic Differentiation (V): Convolution (3)

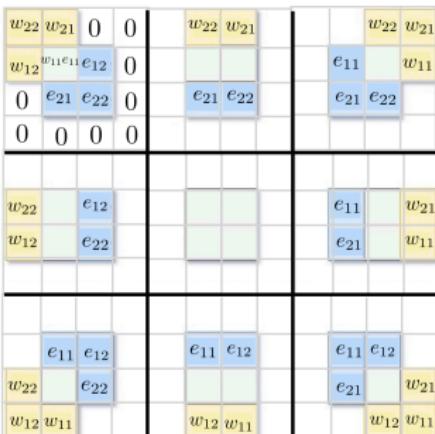
- extend to 2D convolutions
- back-propagation by convolution:

$$\frac{\partial Q}{\partial \mathbf{x}_i} = \sum_{j=1}^k \mathbf{e}_j^{(\emptyset)} * \overleftarrow{\mathbf{w}_{ij}} \quad (i = 1, 2 \dots p)$$

- computing local gradient by convolution:

$$\frac{\partial Q}{\partial \mathbf{w}_{ij}} = \mathbf{x}_i * \mathbf{e}_j \quad (i = 1, 2 \dots p; \ j = 1, 2 \dots k)$$

where  $\mathbf{x}_i \in \mathbb{R}^{d \times d}$  and  $\mathbf{e}_j \in \mathbb{R}^{n \times n}$

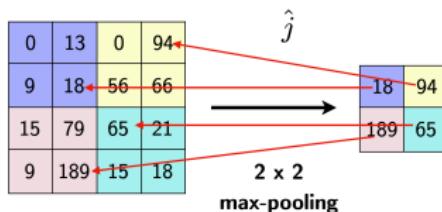


# Automatic Differentiation (VI)

## ■ max-pooling:

- o no parameters  $\Rightarrow$  no local gradients
- o back-propagation:

$$\frac{\partial Q}{\partial x_i} = \begin{cases} \frac{\partial Q}{\partial y_j} & \text{if } i = \hat{j} \\ 0 & \text{otherwise} \end{cases}$$



## ■ batch normalization: $y = BN_{\gamma, \beta}(x)$

- o back-propagation:

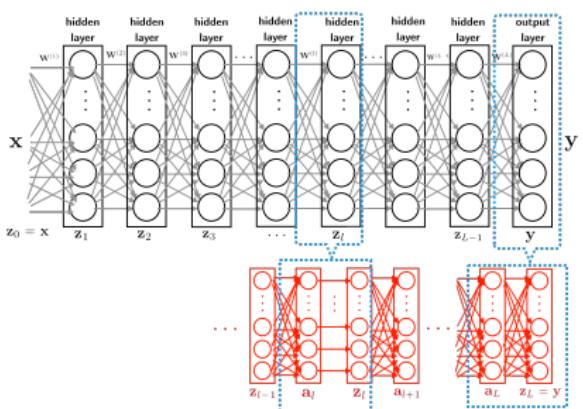
$$\frac{\partial Q}{\partial x^{(m)}} = \frac{M \gamma \odot e^{(m)} - \sum_{k=1}^M \gamma \odot e^{(k)} - \gamma \odot \hat{x}^{(m)} \odot \left( \sum_{k=1}^M e^{(k)} \odot \hat{x}^{(k)} \right)}{M \sqrt{\sigma_B^2(i) + \epsilon}}$$

- o local gradients:

$$\frac{\partial Q}{\partial \gamma} = \sum_{k=1}^M \hat{x}^{(k)} \odot e^{(k)} \quad \frac{\partial Q}{\partial \beta} = \sum_{k=1}^M e^{(k)}$$

# Error Backpropagation Example: Fully-Connected DNNs

- all parameters:  
 $\mathbb{W} = \{\mathbf{W}^{(l)}, \mathbf{b}^{(l)} \mid l = 1, 2 \dots L\}$
- the cross-entropy error:  
$$Q(\mathbb{W}; \mathbf{x}) = -\ln [\mathbf{y}]_r \implies \frac{\partial Q(\mathbb{W}; \mathbf{x})}{\partial \mathbf{y}} = [0 \dots 0 \quad -\frac{1}{y_r} \quad 0 \dots 0]^\top$$
- define error signals  $\mathbf{e}^{(l)} = \frac{\partial Q(\mathbb{W}; \mathbf{x})}{\partial \mathbf{a}_l}$  for all  $l = L, \dots, 2, 1$
- apply AD to the softmax, nonlinear activation and full connection modules to back-propagate error signals



# Error Backpropagation Example: Fully-Connected DNNs

backward pass of fully-connected DNNs

for the cross-entropy error of any input-output pair  $(\mathbf{x}, \mathbf{r})$

- 1 for the output layer  $L$ :

$$\mathbf{e}^{(L)} = [y_1 \ y_2 \ \cdots \ y_r - 1 \ \cdots \ y_n]^\top$$

- 2 for each hidden layer  $l = L - 1, \dots, 2, 1$ :

$$\mathbf{e}^{(l)} = \left( (\mathbf{W}^{(l+1)})^\top \mathbf{e}^{(l+1)} \right) \odot H(\mathbf{z}_l)$$

- 3 for all layers  $l = L, \dots, 2, 1$ :

$$\frac{\partial Q(\mathbb{W}; \mathbf{x})}{\partial \mathbf{W}^{(l)}} = \mathbf{e}^{(l)} (\mathbf{z}_{l-1})^\top$$

$$\frac{\partial Q(\mathbb{W}; \mathbf{x})}{\partial \mathbf{b}^{(l)}} = \mathbf{e}^{(l)}$$

where  $\mathbf{y}$  and  $\mathbf{z}_l$  ( $l = 0, 1, \dots, L - 1$ ) are computed in the forward pass

# Heuristics and Tricks for Optimization

- Hyperparameters
- Optimization Method: ADAM
- Regularization
- Fine-tuning Tricks

# Hyperparameters of Learning Neural Networks

- initial parameters
- epoch number
- mini-batch size
- learning rate
  - a good initial learning rate  $\eta_0$
  - an annealing schedule to adjust  $\eta_t \rightarrow \eta_{t+1}$
  - call for some self-adjusting mechanisms, e.g. Adagrad, Adadelta, ADAM, AdaMax, etc.

# Optimization method: ADAM

## ADAM to learn neural networks

randomly initialize  $\mathbb{W}^{(0)}$ , and set  $\eta$ ,  $t = 0$ ,  $n = 0$  and  $\mathbf{u}_0 = \mathbf{v}_0 = 0$

**while** not converged **do**

    randomly shuffle training data into mini-batches

**for** each mini-batch  $B$  **do**

**for** each  $\mathbf{x} \in B$  **do**

            compute  $\frac{\partial Q(\mathbb{W}^{(n)}; \mathbf{x})}{\partial \mathbb{W}}$

**end for**

$\mathbf{g}_n = \frac{1}{|B|} \sum_{\mathbf{x} \in B} \frac{\partial Q(\mathbb{W}^{(n)}; \mathbf{x})}{\partial \mathbb{W}}$

$\mathbf{u}_{n+1} = \alpha \mathbf{u}_n + (1 - \alpha) \mathbf{g}_n$  and  $\mathbf{v}_{n+1} = \beta \mathbf{v}_n + (1 - \beta) \mathbf{g}_n \odot \mathbf{g}_n$

$\hat{\mathbf{u}}_{n+1} = \frac{\mathbf{u}_{n+1}}{1 - \alpha^{n+1}}$  and  $\hat{\mathbf{v}}_{n+1} = \frac{\mathbf{v}_{n+1}}{1 - \beta^{n+1}}$

        update model:  $\mathbb{W}^{(n+1)} = \mathbb{W}^{(n)} - \eta \cdot \hat{\mathbf{u}}_{n+1} \odot ((\hat{\mathbf{v}}_{n+1} + \epsilon^2)^{-\frac{1}{2}})$

$n = n + 1$

**end for**

$t = t + 1$

**end while**



# Self-adjusting Mechanism in ADAM

- use exponential average to accumulate 1st-order and 2nd-order moments ( $\mathbf{u}_n$  and  $\mathbf{v}_n$ ) of the gradient ( $\mathbf{g}_n$ )
- normalize to yield unbiased estimates:

$$\mathbb{E}[\hat{\mathbf{u}}_{n+1}(i)] = \mathbb{E}[\mathbf{g}_n(i)] \quad \mathbb{E}[\hat{\mathbf{v}}_{n+1}(i)] = \mathbb{E}[\mathbf{g}_n^2(i)]$$

- model update formula:

$$\mathbb{W}_i^{(n+1)} = \mathbb{W}_i^{(n)} - \eta \frac{\hat{\mathbf{u}}_{n+1}(i)}{\sqrt{\hat{\mathbf{v}}_{n+1}(i) + \epsilon^2}}$$

- self-adjusting model updates  $\Delta \mathbb{W}_i^{(n)}$ :

$$\|\Delta \mathbb{W}_i^{(n)}\|^2 \simeq \eta^2 \frac{(\mathbb{E}[\hat{\mathbf{u}}_{n+1}(i)])^2}{\mathbb{E}[\hat{\mathbf{v}}_{n+1}(i)]} = \frac{\eta^2 (\mathbb{E}[\mathbf{g}_n(i)])^2}{(\mathbb{E}[\mathbf{g}_n(i)])^2 + \text{var}[\mathbf{g}_n(i)]}$$

# Regularization in Neural Networks

- **weight decay:** use  $L_2$  norm regularization

$$Q(\mathbb{W}) + \frac{\lambda}{2} \cdot \|\mathbb{W}\|^2 \implies \mathbb{W}^{(n+1)} = \mathbb{W}^{(n)} - \eta \frac{\partial Q(\mathbb{W}^{(n)})}{\partial \mathbb{W}} - \lambda \cdot \mathbb{W}^{(n)}$$

- **weight normalization:** normalize weight vectors to facilitate optimization

1. tied-scalar reparameterization:

$$\mathbf{w} = \gamma \cdot \mathbf{v} \quad \text{s.t.} \quad \|\mathbf{v}\| \leq 1$$

2. normalizing reparameterization:

$$\mathbf{w} = \frac{\gamma}{\|\mathbf{v}\|} \mathbf{v}$$

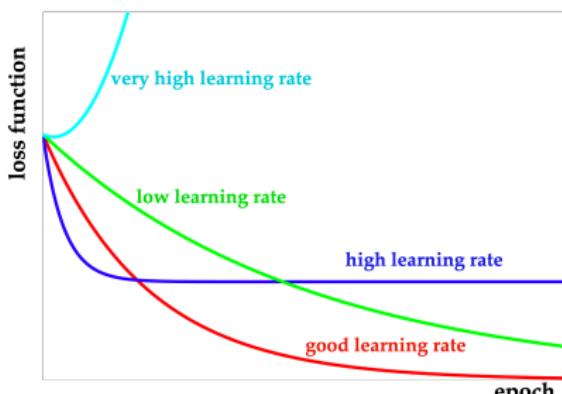
- **dropout**

- **data augmentation**

# Fine-tuning Tricks

critical to monitor three learning curves:

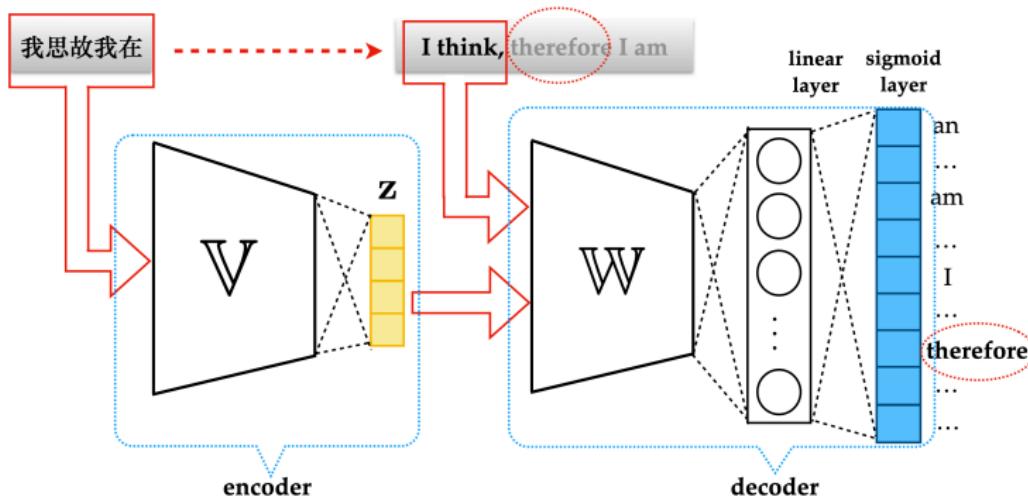
- the objective function (a.k.a. loss function)
- performance on training data
- performance on development data



# End-to-End Learning

- **end-to-end learning:** train a single model to map directly from raw data to final targets
- neural networks are suitable for end-to-end learning
  - flexible architectures to accommodate a variety of raw data
  - powerful enough to approximate potentially complex mapping
  - arrange output structures to generate real data, e.g. *deconvolution* layers for images, *WaveNet* for audio/speech
- the popular **encoder-decoder** structure
- **sequence-to-sequence learning:** learn deep neural networks to map from one input sequence to an output sequence
  - suitable for many NLP tasks, e.g. machine translation, question-answering, etc.

# Sequence-to-Sequence Learning



- encoder and decoder are powerful neural networks that can handle sequences, e.g. RNNs, LSTMs, or transformers